

ML & Neural Approaches for Information Retrieval (Part 1)

Andrew Yates
Assistant Professor, UvA

[@andrewyates](https://twitter.com/andrewyates)

July 20, 2022
The 13th European Summer School in Information Retrieval



UNIVERSITEIT VAN AMSTERDAM

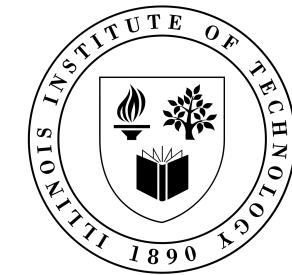
About me

Assistant Professor at University of Amsterdam (2021-)

- Information retrieval, neural networks, natural language processing, (consumer) medical domain, personal knowledge graphs, ...

Previously

- Postdoc then Senior Researcher (2018-2021) at Max Planck Institute for Informatics
- PhD (2016) in Computer Science from Georgetown University, Washington, DC
- BSc in Computer Science from Illinois Institute of Technology, Chicago



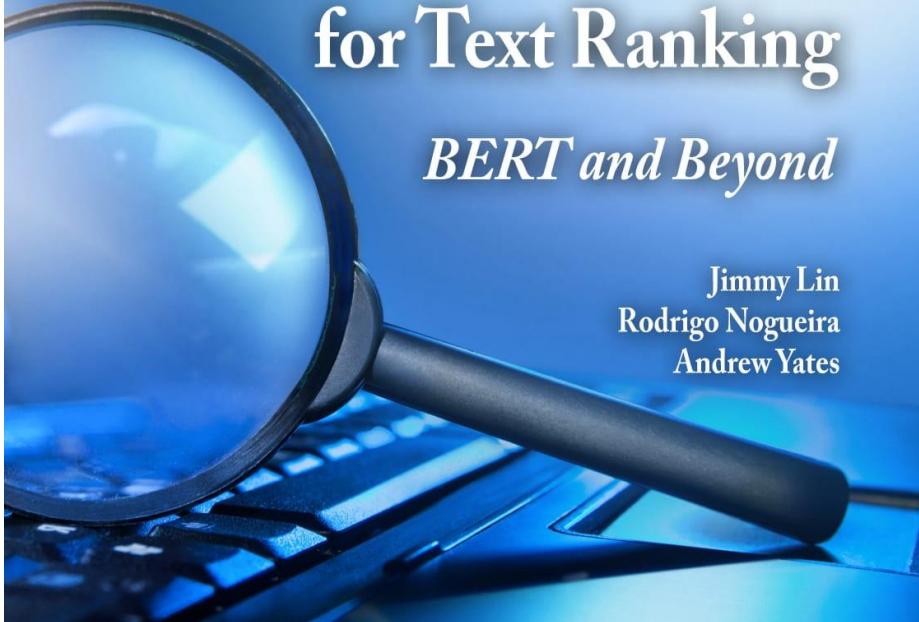


MORGAN & CLAYPOOL PUBLISHERS

Pretrained Transformers for Text Ranking

BERT and Beyond

Jimmy Lin
Rodrigo Nogueira
Andrew Yates



SYNTHESIS LECTURES
HUMAN LANGUAGE TECHNOLOGIES

Graeme Hirst, Series Editor

Collaborators: Siddhant Arora, Arman Cohan, Jeffrey Dalton, Doug Downey, Sergey Feldman, Nazli Goharian, Kevin Martin Jose, Jimmy Lin, Sean MacAvaney, Thong Nguyen, Rodrigo Nogueira, Wei Yang, Xinyu Zhang, ...



**Capreolus: Toolkit for
Neural Ad Hoc Retrieval**

DiffIR

when did rock n roll begin?

query_id 940547

text when did rock n roll begin?

trec-dl-2020/p_bm25

trec-dl-2020/p_bm25rm3

1 ↗8 Rel: 0 Score: 15.0555

text: ...with **rock 'n' roll** didn't **begin** and end with the TURTLES legacy; there is much more to the story. Mark was part of the LA's in-crowd in the '60s, hanging out with some of the biggest names in **rock 'n' r...**

2 ↗6 Rel: 0 Score: 14.4202

text: ...g of **Rock 'n' Roll** at Elvis' home, Graceland. The full Graceland experience will take you from Elvis' humble beginnings through his rise to superstardom. See how a **rock 'n' roll** legend lived and relax...

3 ↗12 Rel: 0 Score: 13.9817

text: ...s of **rock 'n' roll** must also be attributed to the influence of black culture and music. The blues had an illegitimate baby and we named it **rock 'n' roll**, Little Richard, one of the originators of **rock...**

ir_datasets: Catalog

✓: Data available as automatic download

⚠: Data available from a third party

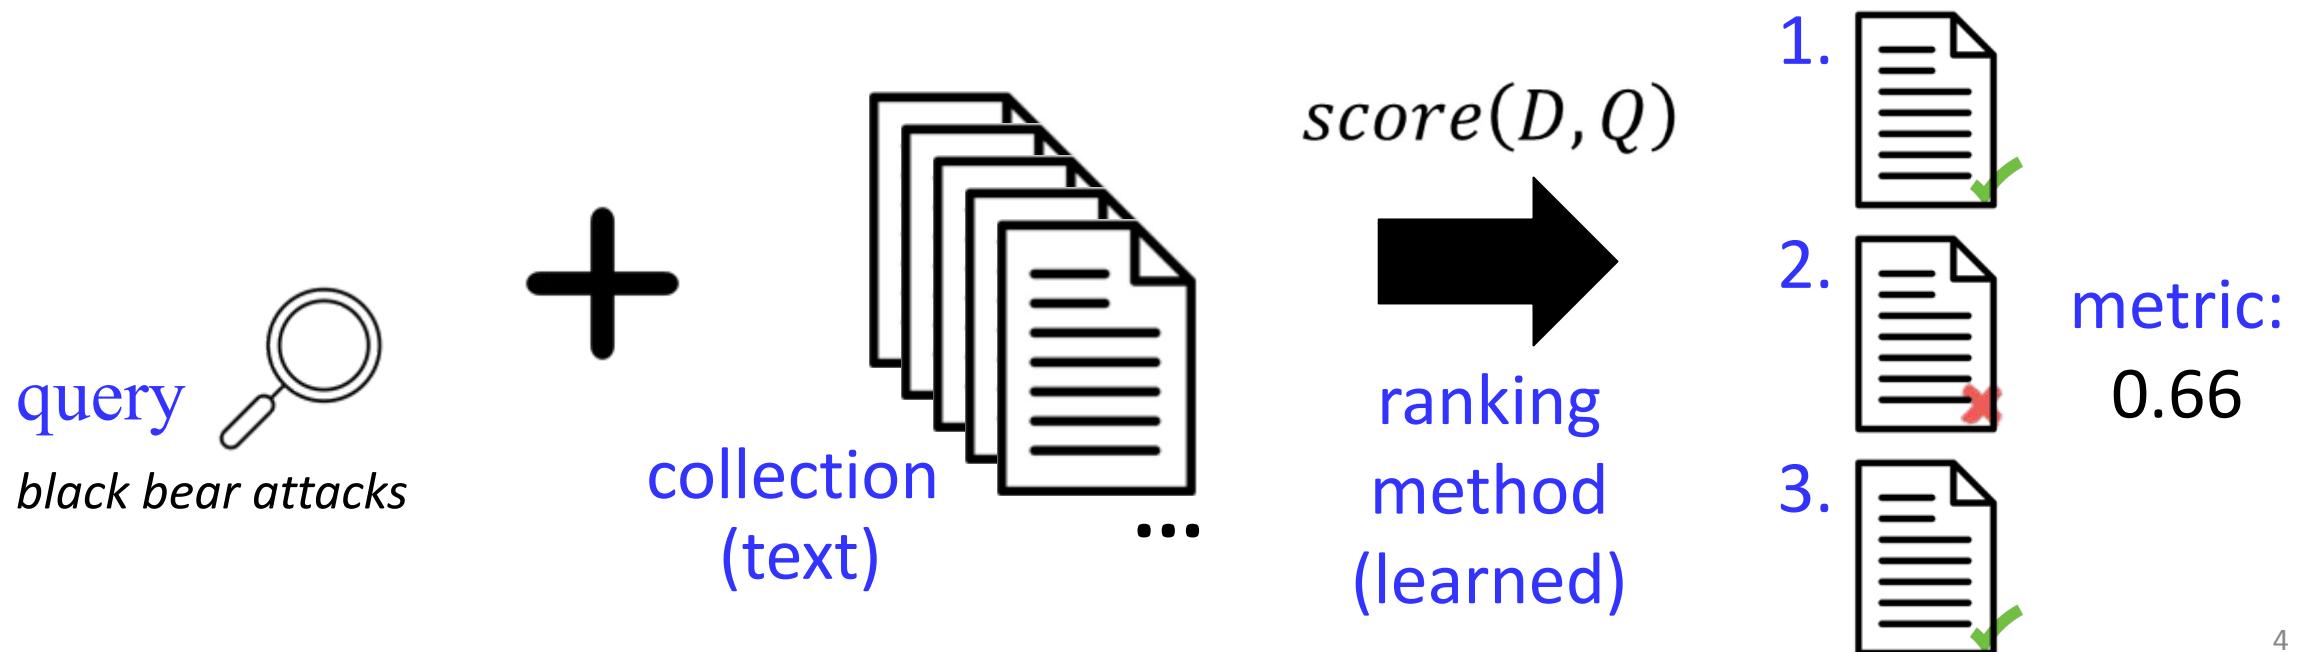
↑: Data inherited from a parent dataset (highlights which one on hover)

Dataset	docs	queries	qrels
antique	✓		
antique/test	↑	✓	✓
antique/test/non-offensive	↑	✓	✓
antique/train	↑	✓	✓
antique/train/split200-train	↑	✓	✓
antique/train/split200-valid	↑	✓	✓
aol-ia	⚠	✓	✓

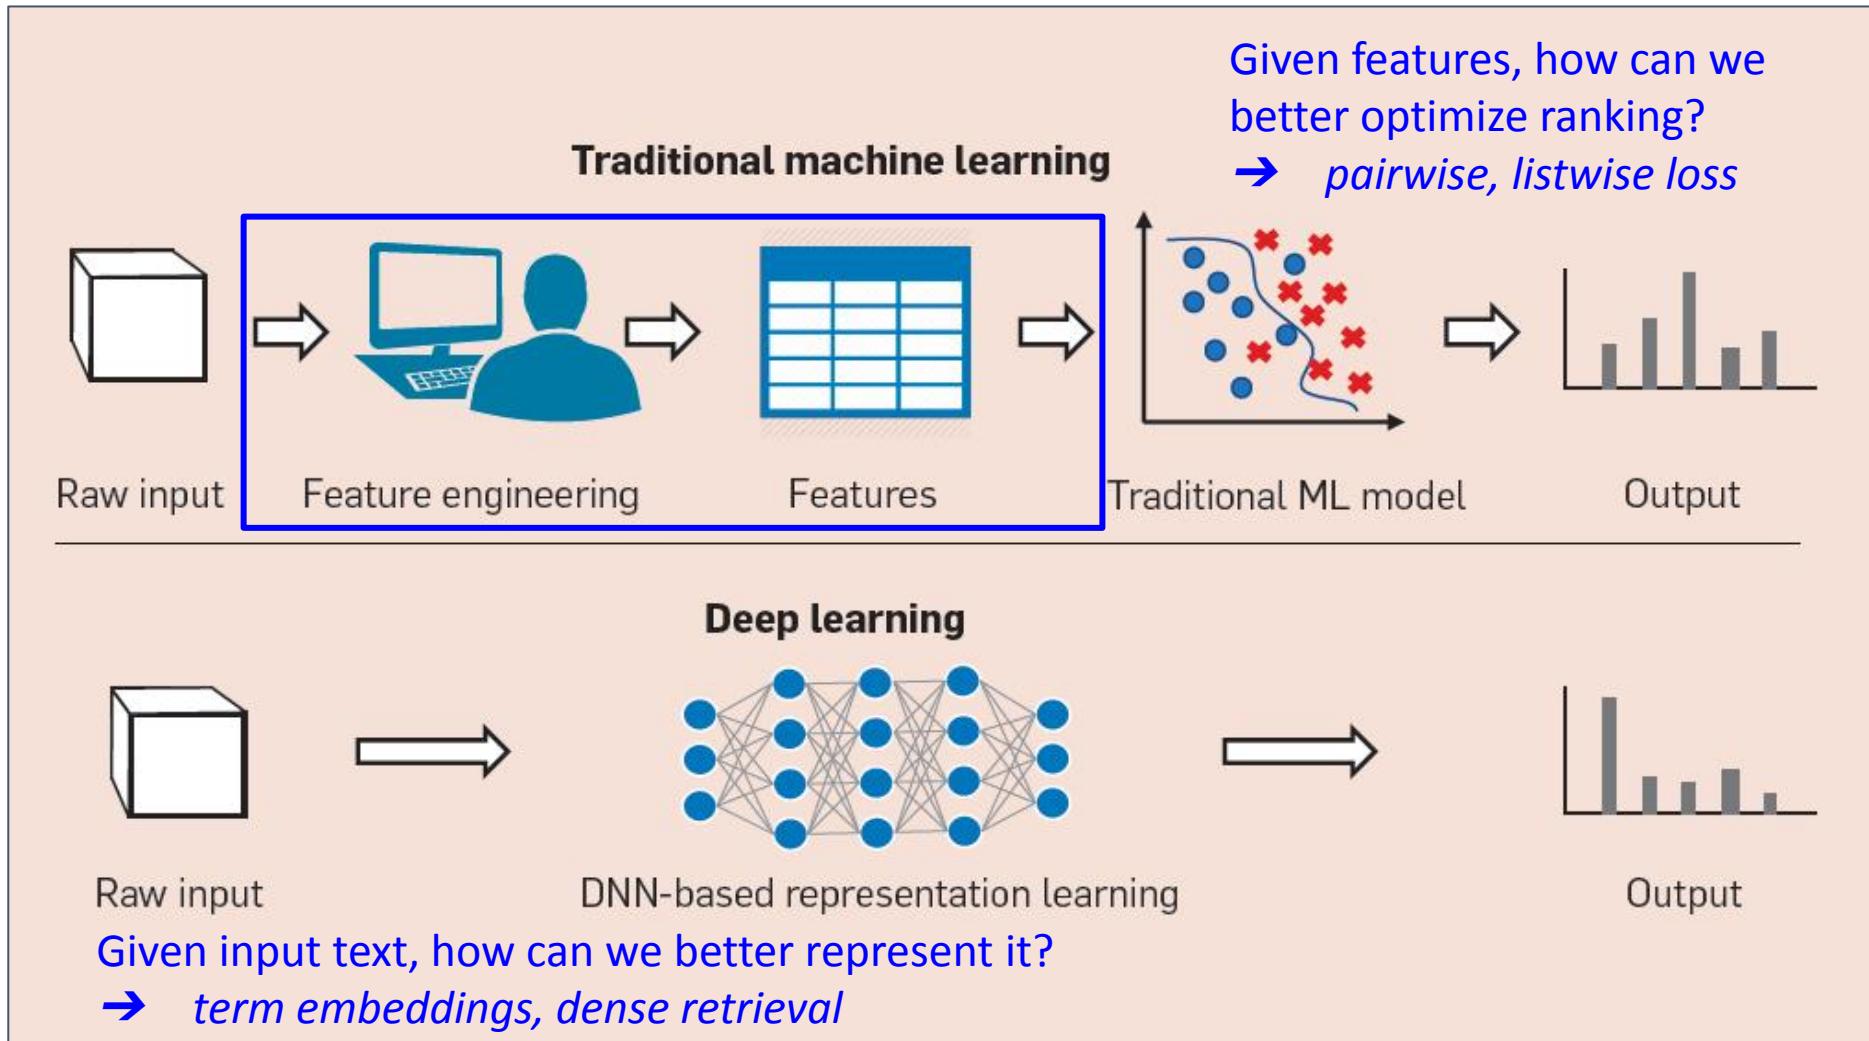
Introduction

Methods that use relevance judgments to **learn** how to rank results

- **ML** approaches **learn to rank** based on hand-crafted features (e.g., BM25 score between Q and D, spam score for D, etc)
- **Neural** approaches learn improved **text representations** for ranking



Introduction



Advantage #1: less handcrafting

$$score(D, Q) = \sum_{i=1}^{|Q|} IDF(q_i) * TF(q_i, D)$$

Inverse Document
Frequency of query term
(collection stat)

Term Frequency of
 q in document D
(document stat)

$N = \# \text{ docs}$ and $n_i = \# \text{ docs with term } q_i$

IDF defined as:

- $\log N/n_i$
- $\log (1 + N/n_i)$
- $\log (N - n_i)/n_i$
- ...

$$score(D, Q) = \sum_{i=1}^{|Q|} IDF(q_i) * TF(q_i, D)$$

Inverse Document
Frequency of query term

Term Frequency of
 q in document D

TF defined as:

- $TF(q_i, D)$
- $1 + \log TF(q_i, D)$
- $0.5 + 0.5 [TF(q_i, D) / \max TF(D)]$
- ...

$$score(D, Q) = \sum_{i=1}^{|Q|} IDF(q_i) * TF(q_i, D)$$

Inverse Document
Frequency of query term

Term Frequency of
 q in document D

Advantage #2: soft matching of terms

$$score(D, Q) = \sum_{i=1}^{|Q|} IDF(q_i) * TF(q_i, D)$$



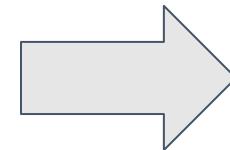
Exact matching of
Q and D terms

Non-neural approaches: translation models, topic models like pLSI, ...

Advantage #2: soft matching of terms

document ranking
text ranking
document retrieval
Web search

Match?



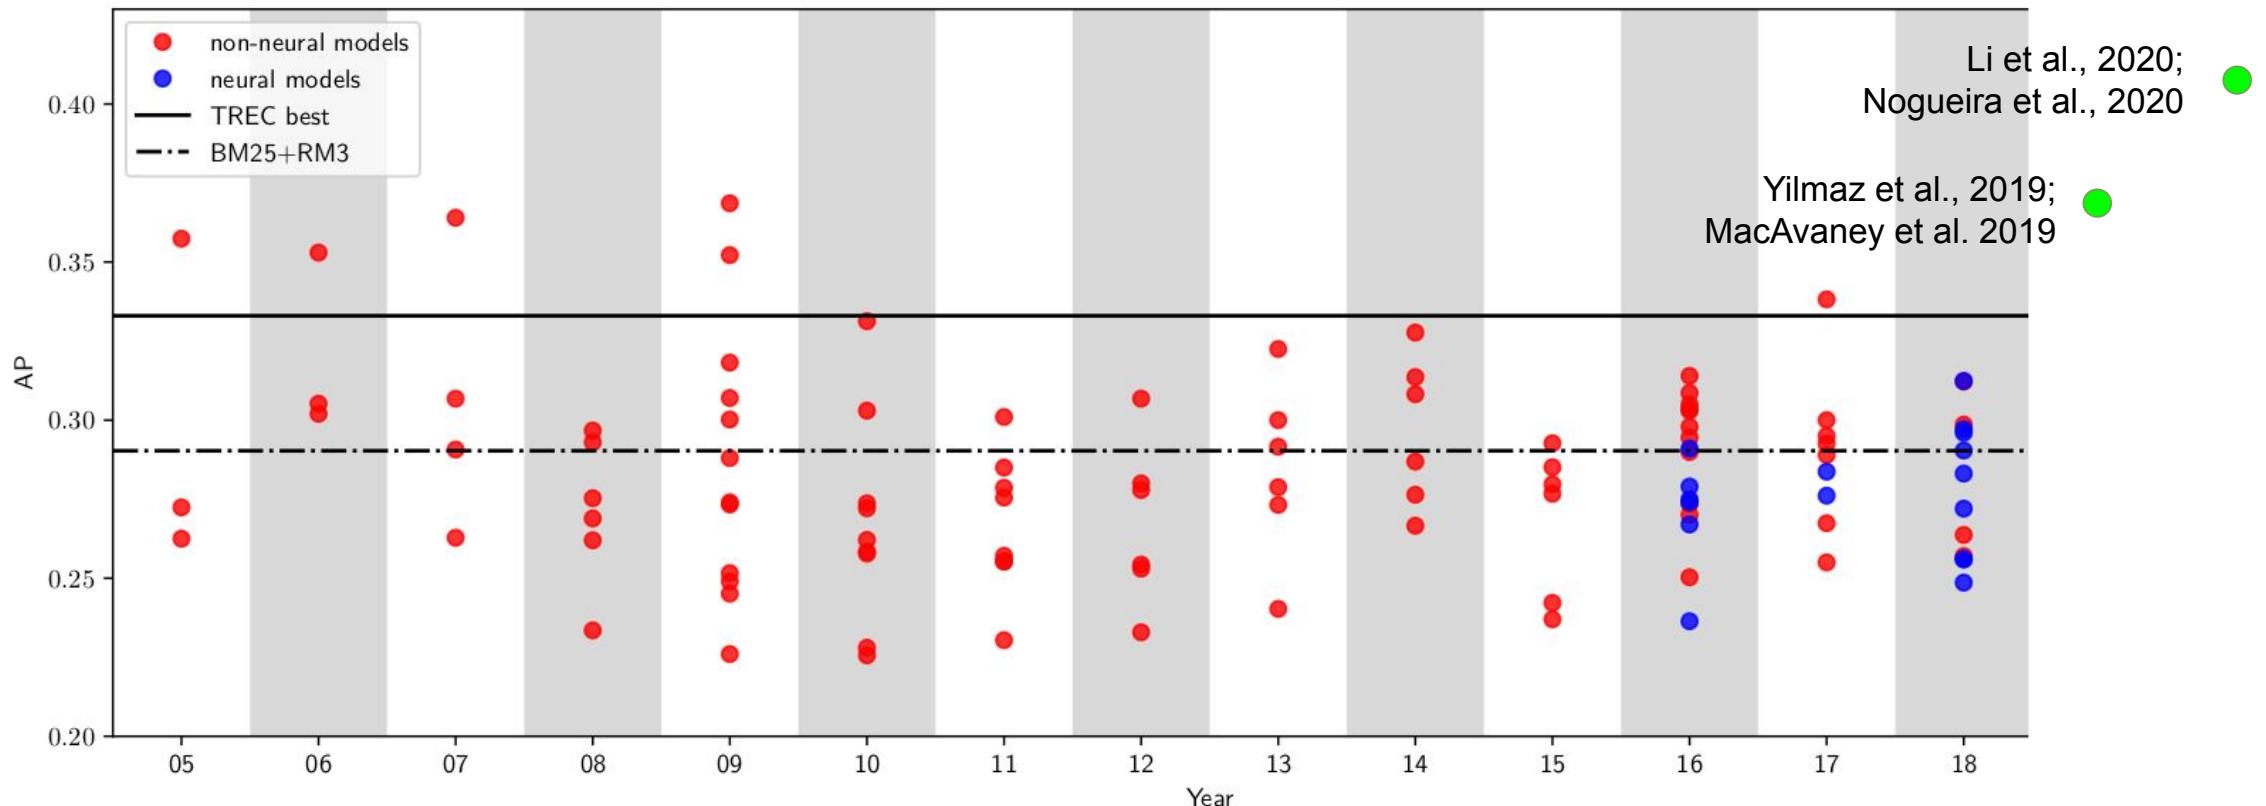
query, search,
ranking, retrieval, ...

Enrich query or document representations to move beyond exact matching

- Unsupervised: pseudo-relevance feedback
- Neural: embeddings, document expansion, ...

Search quality improvements

Robust04 Dataset (news articles)



Search quality improvements

MS MARCO Passage Dataset (Web pages)

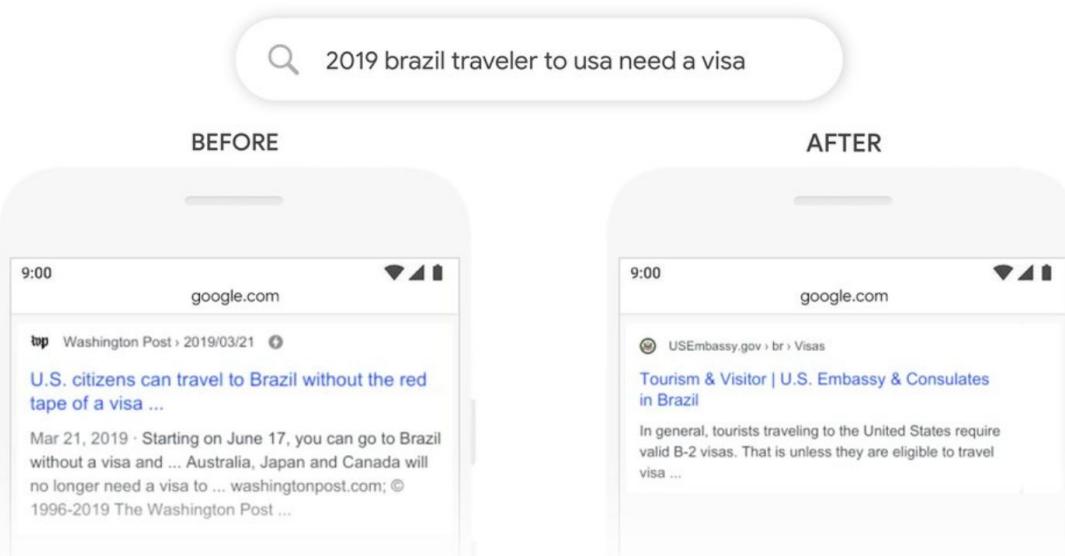
Rank	Model	Submission Date	MRR@10 On Eval	MRR@10 On Dev
1	BERT + Small Training Rodrigo Nogueira and Kyunghyun Cho - New York University [Nogueira, et al. '19] and [Code]	January 7th, 2019	35.87	~8 points!
2	IRNet (Deep CNN/IR Hybrid Network) Dave DeBarr, Navendu Jain, Robert Sim, Justin Wang, Nirupama Chandrasekaran – Microsoft	January 2nd, 2019	28.06	
3	Neural Kernel Match IR (Conv-KNRM) (1)Yifan Qiao, (2)Chenyan Xiong, (3)Zhenghao Liu, (4)Zhiyuan Liu-Tsinghua University(1, 3, 4); Microsoft Research AI(2) [Dai et al. '18]	Novmeber 28th, 2018	27.12	29.02
4	Neural Kernel Match IR (KNRM) ((1)Yifan Qiao, (2)Chenyan Xiong, (3)Zhenghao Liu, (4)Zhiyuan Liu-Tsinghua University(1, 3, 4); Microsoft Research AI(2) [Xiong et al. '17]	December 10th, 2018	19.82	21.84
5	Feature-based LeToR: simple-feature based RankSVM (1)Yifan Qiao, (2)Chenyan Xiong, (3)Zhenghao Liu, (4)Zhiyuan Liu-Tsinghua University(1, 3, 4); Microsoft Research AI(2)	December 10th, 2018	19.05	19.47
6	BM25	Novmeber 1st, 2018	16.49	16.70

19 point improvement: BM25 vs. best neural

8 point improvement: neural pre-LLM vs. LLM

Search quality improvements: industry

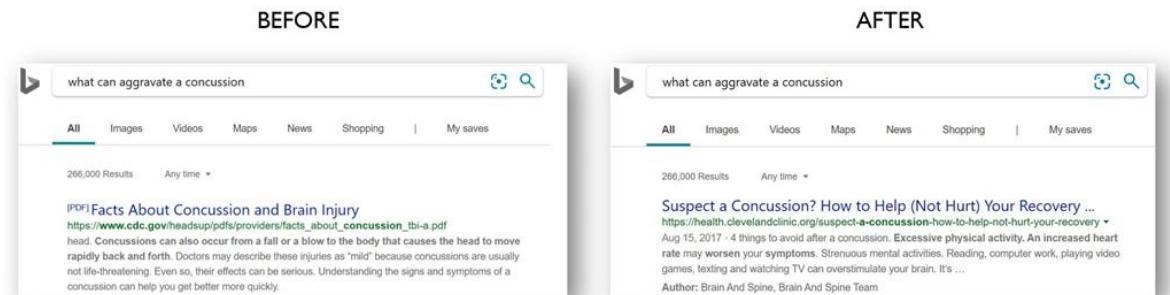
Google Search



*We're making a significant improvement to how we understand queries, representing the **biggest leap forward in the past five years**, and **one of the biggest leaps forward in the history of Search**.*

[source](#)

MS Bing



Starting from April of this year (2019), we used large transformer models to deliver the largest quality improvements to our Bing customers in the past year.

[source](#)

Outline

- Introduction
- **Ranking with soft matching**
- Transformers & contextualized embeddings
- Transformer approaches for ranking
 - Re-ranking with cross-encoders
 - Learned sparse retrieval
 - Document expansion
- Conclusion

From exact matching to soft matching

Key point: neural methods replace *exact* matching with *soft* matching

With traditional methods, soft matching is possible (but it's less effective)

- *Stemming* handles matches like singular vs. plural and different tenses
swam, swimming, swim, swims replaced with *swim*
- Query expansion approaches handle context by adding related terms to query
Query contains *bank*, referring to turning a plane
Add related terms to query: *flight, airplane, ailerons, spoilers, ...*

Soft matching via embeddings

	living being	feline	human	gender	royalty	verb	plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8

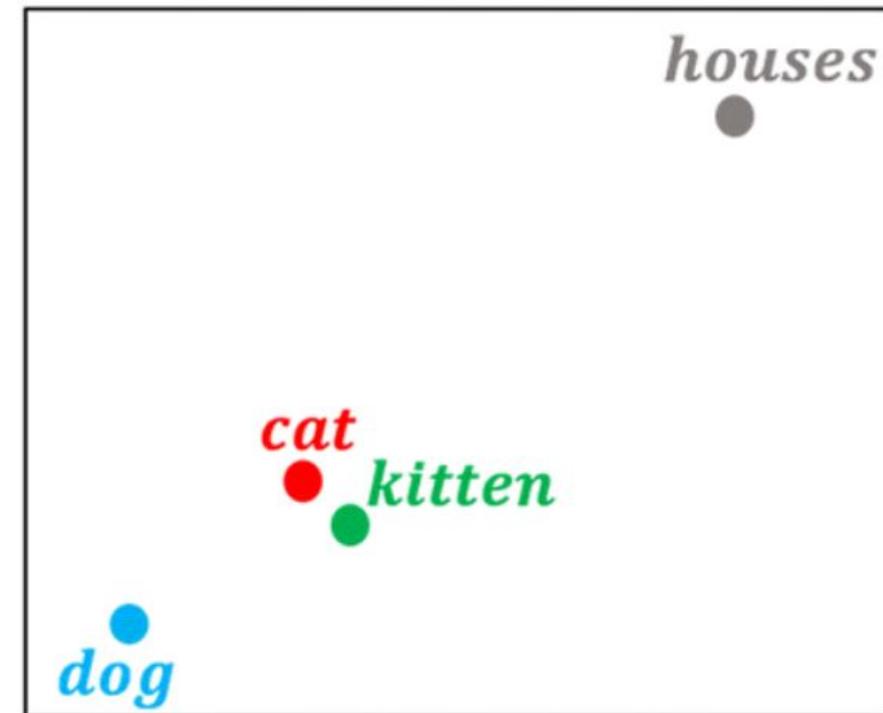
Neural approaches
represent words
as learned *vectors*

Soft matching via embeddings

Vector “*embedding*” allows comparisons of different terms

	living	being	feline	human	gender	royalty	verb	plural
<i>cat</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2	
<i>kitten</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1	
<i>dog</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3	
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8	

Dimensionality
reduction of
word
embeddings
from 7D to 2D



Ranking with embeddings

Score can be produced by placing document terms in *similarity buckets*, then computing relevance based on the size of each bucket

Query Terms		sim = 1.0	0.6 < sim < 1.0	sim <= 0.6
	curbing	4 terms	3 terms	20 terms
	population	1 term	8 terms	40 terms
	growth	2 terms	2 terms	15 terms
Similarity Buckets				

Weights: a, b, c

$$\text{Score}(\text{curbing}, D) = 4*a + 3*b + 20*c$$

$$\text{Score}(Q, D) = \text{Score}(\text{curbing}, D) + \text{Score}(\text{population}, D) + \text{Score}(\text{growth}, D)$$

Deep Relevance Matching Model (DRMM)

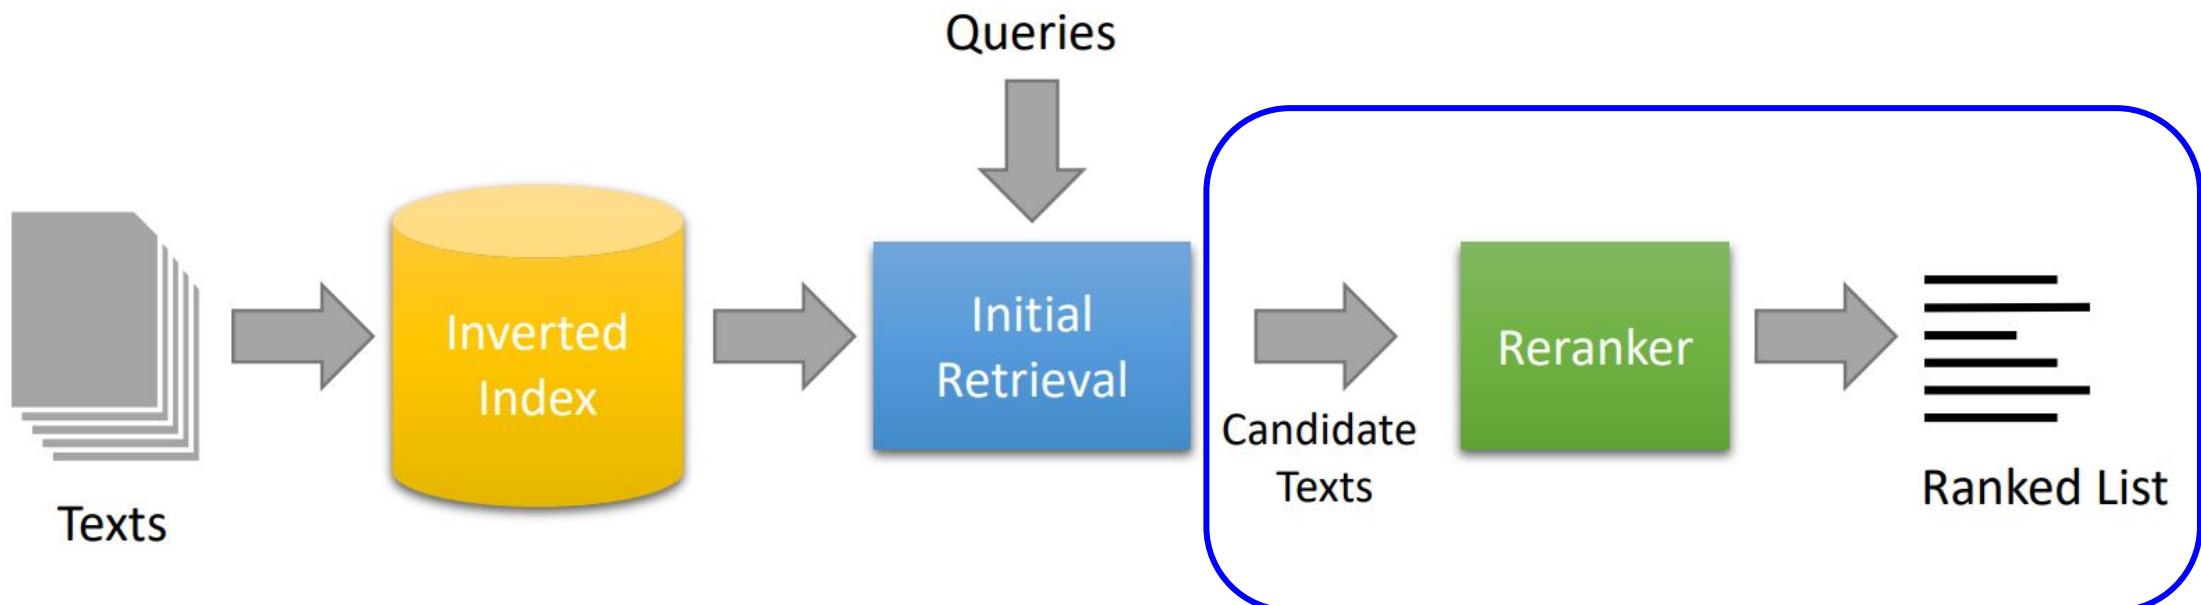
Input: query Q and a document D to score

Scoring procedure:

Load document & embeddings into memory, then compute similarity with query

1. Load document D from the forward index, representing it as a list of term IDs
2. Use term IDs to index into embedding matrix, representing D as a list of embeddings
3. Compute histogram $h(t)$ for each query term,
capturing the cosine similarities between the embeddings of t and each doc term
4. Compute term score $z(h(t))$ using a feedforward network
5. Compute weight $g(t)$ for each query term (IDF or using embedding)
6. Compute relevance score as summation of $z(h(t)) * g(t)$ over all query terms

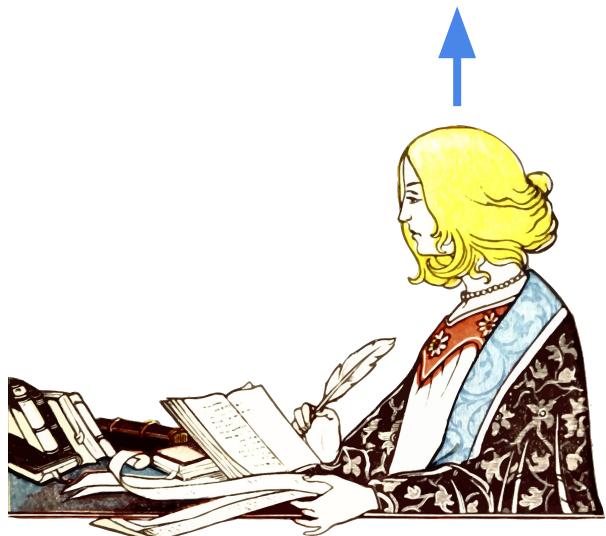
Reranking with DRMM



Training DRMM

Loss:

$$L = - \sum_{j \in J_{\text{pos}}} \log(s_j) - \sum_{j \in J_{\text{neg}}} \log(1 - s_j)$$



Humans



BM25

Outline

- Introduction
- Ranking with soft matching
- **Transformers & contextualized embeddings**
- Transformer approaches for ranking
 - Re-ranking with cross-encoders
 - Learned sparse retrieval
 - Document expansion
- Conclusion

Contextualized embeddings

Idea: a word's representation should vary with context

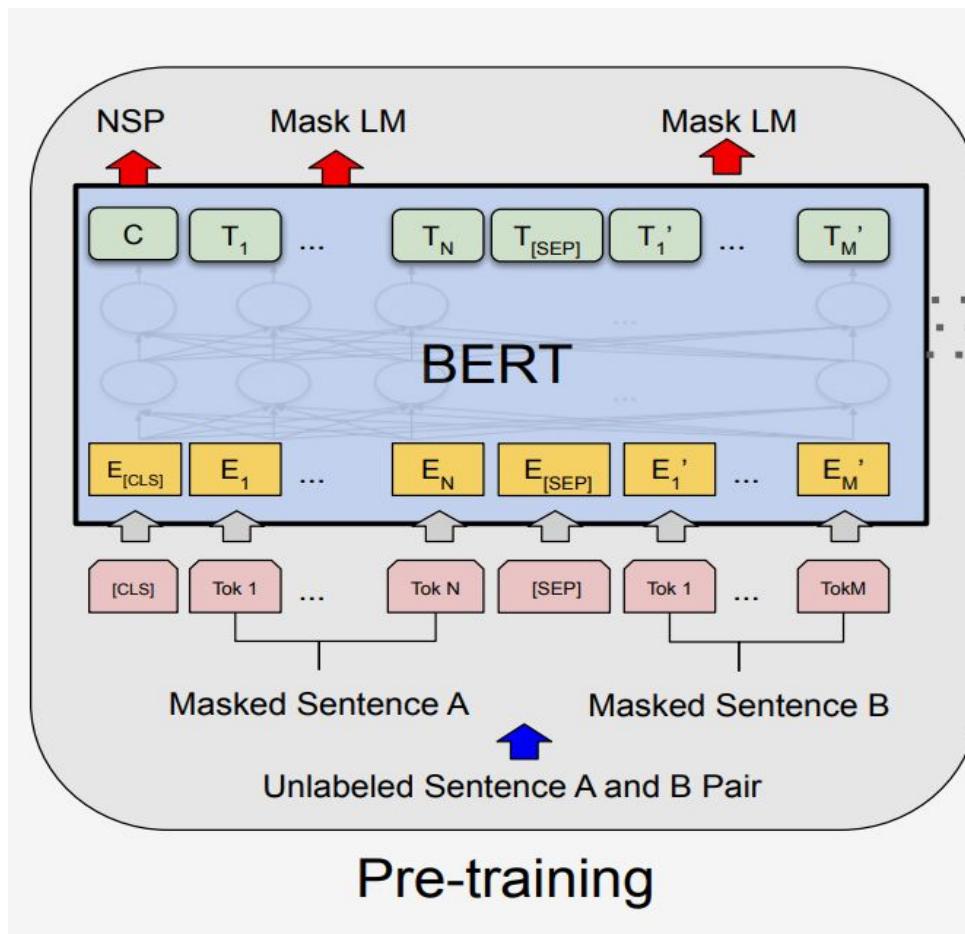
- Generate embeddings based on given text (e.g., query, document, sentence)
- If the context is unique, so is the embedding

In contrast with word2vec, GloVe, FastText, etc, which

- learn one static embedding per word
- learn embeddings based on co-occurrences of word pairs
 - e.g., (pet, dog) more likely than (pet, alligator) ... than (pet, taxes)

Canonical approach: BERT

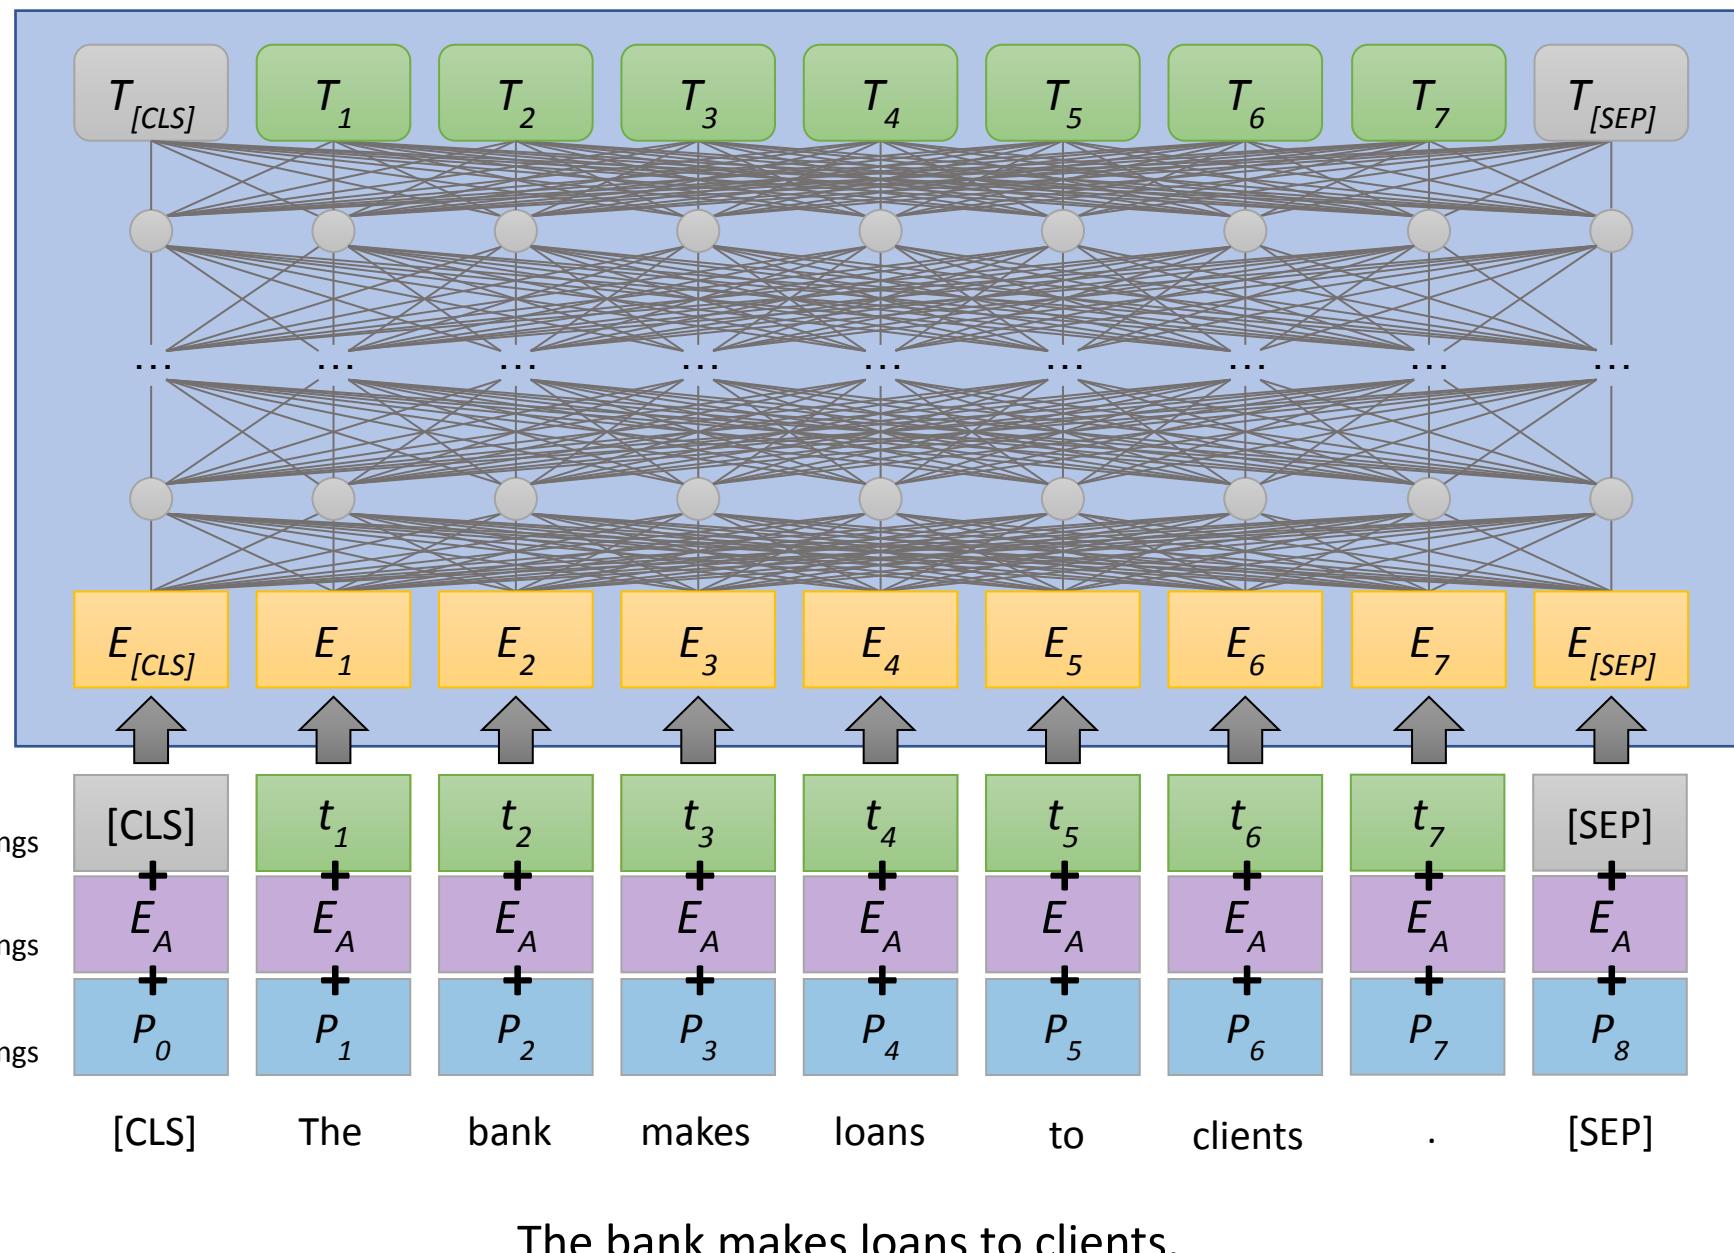
What is BERT?



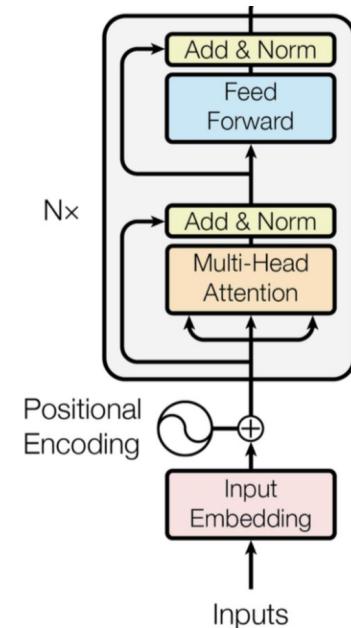
Self-supervised: ∞ training data

BERT

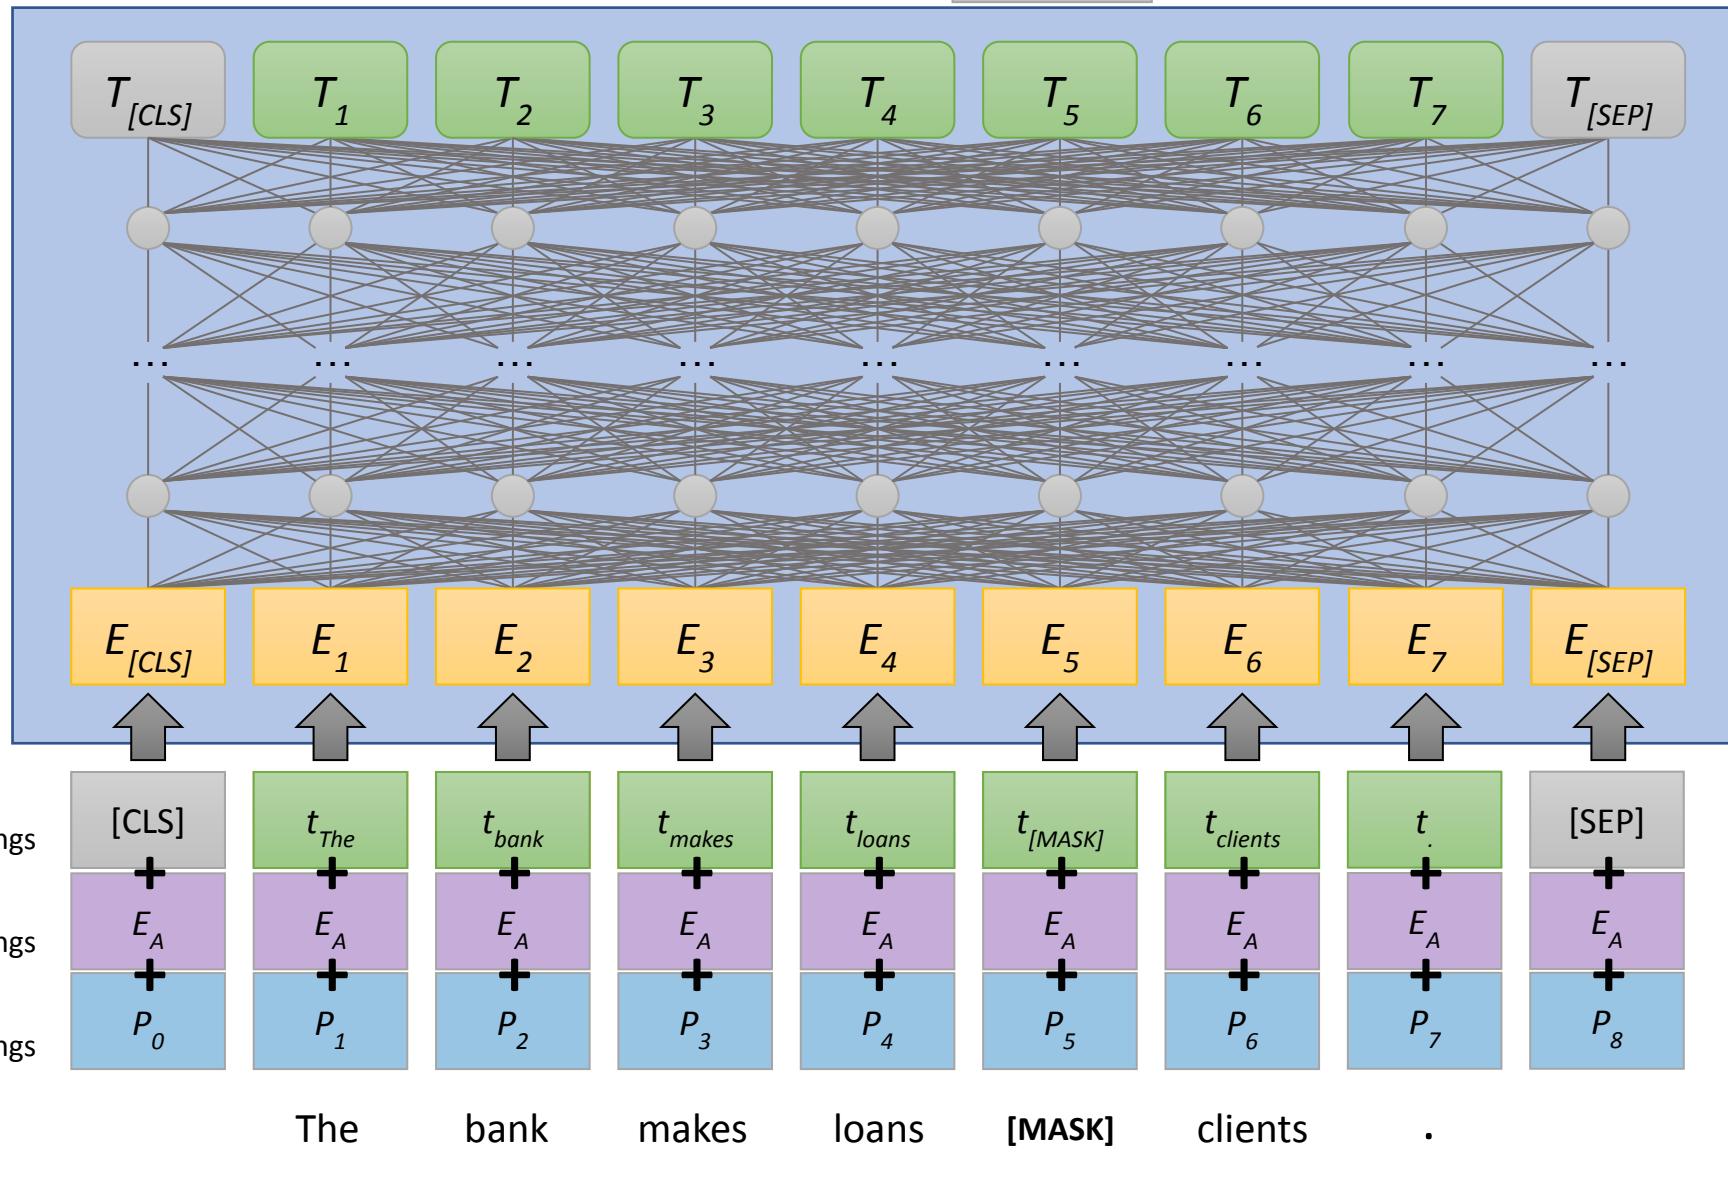
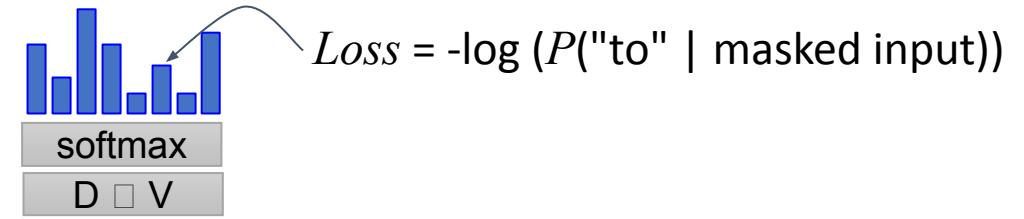
string →
sequence of
vectors



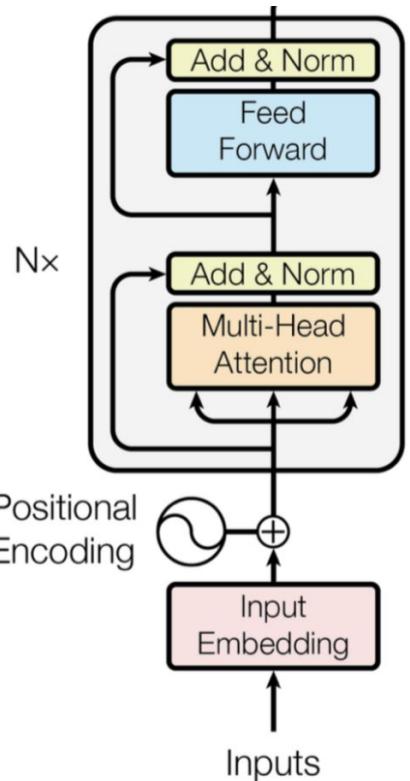
Transformer
Encoder



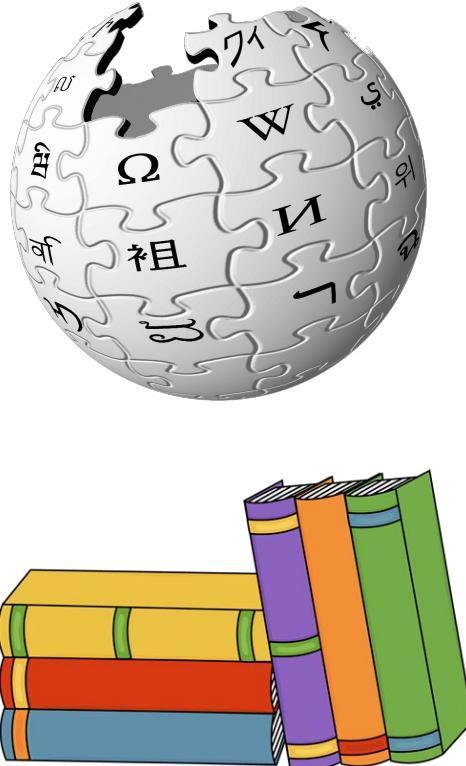
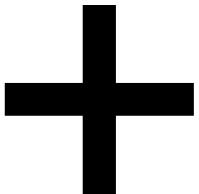
Pretraining: Masked Language Modeling



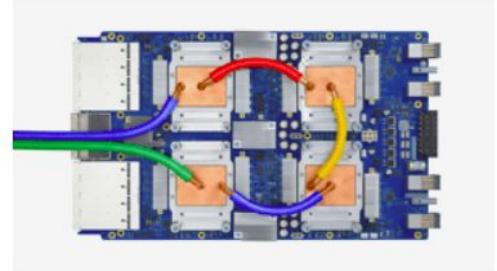
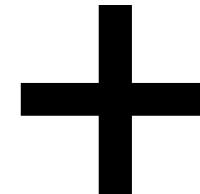
Pretraining ingredients



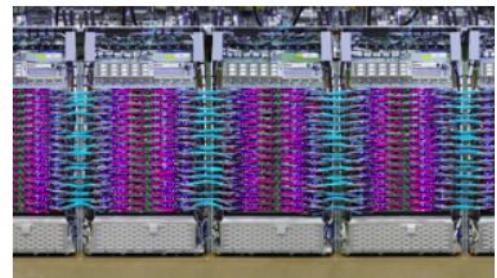
Lots of parameters
(stack of transformers)



Lots of text



Cloud TPU v3
420 teraflops
128 GB HBM



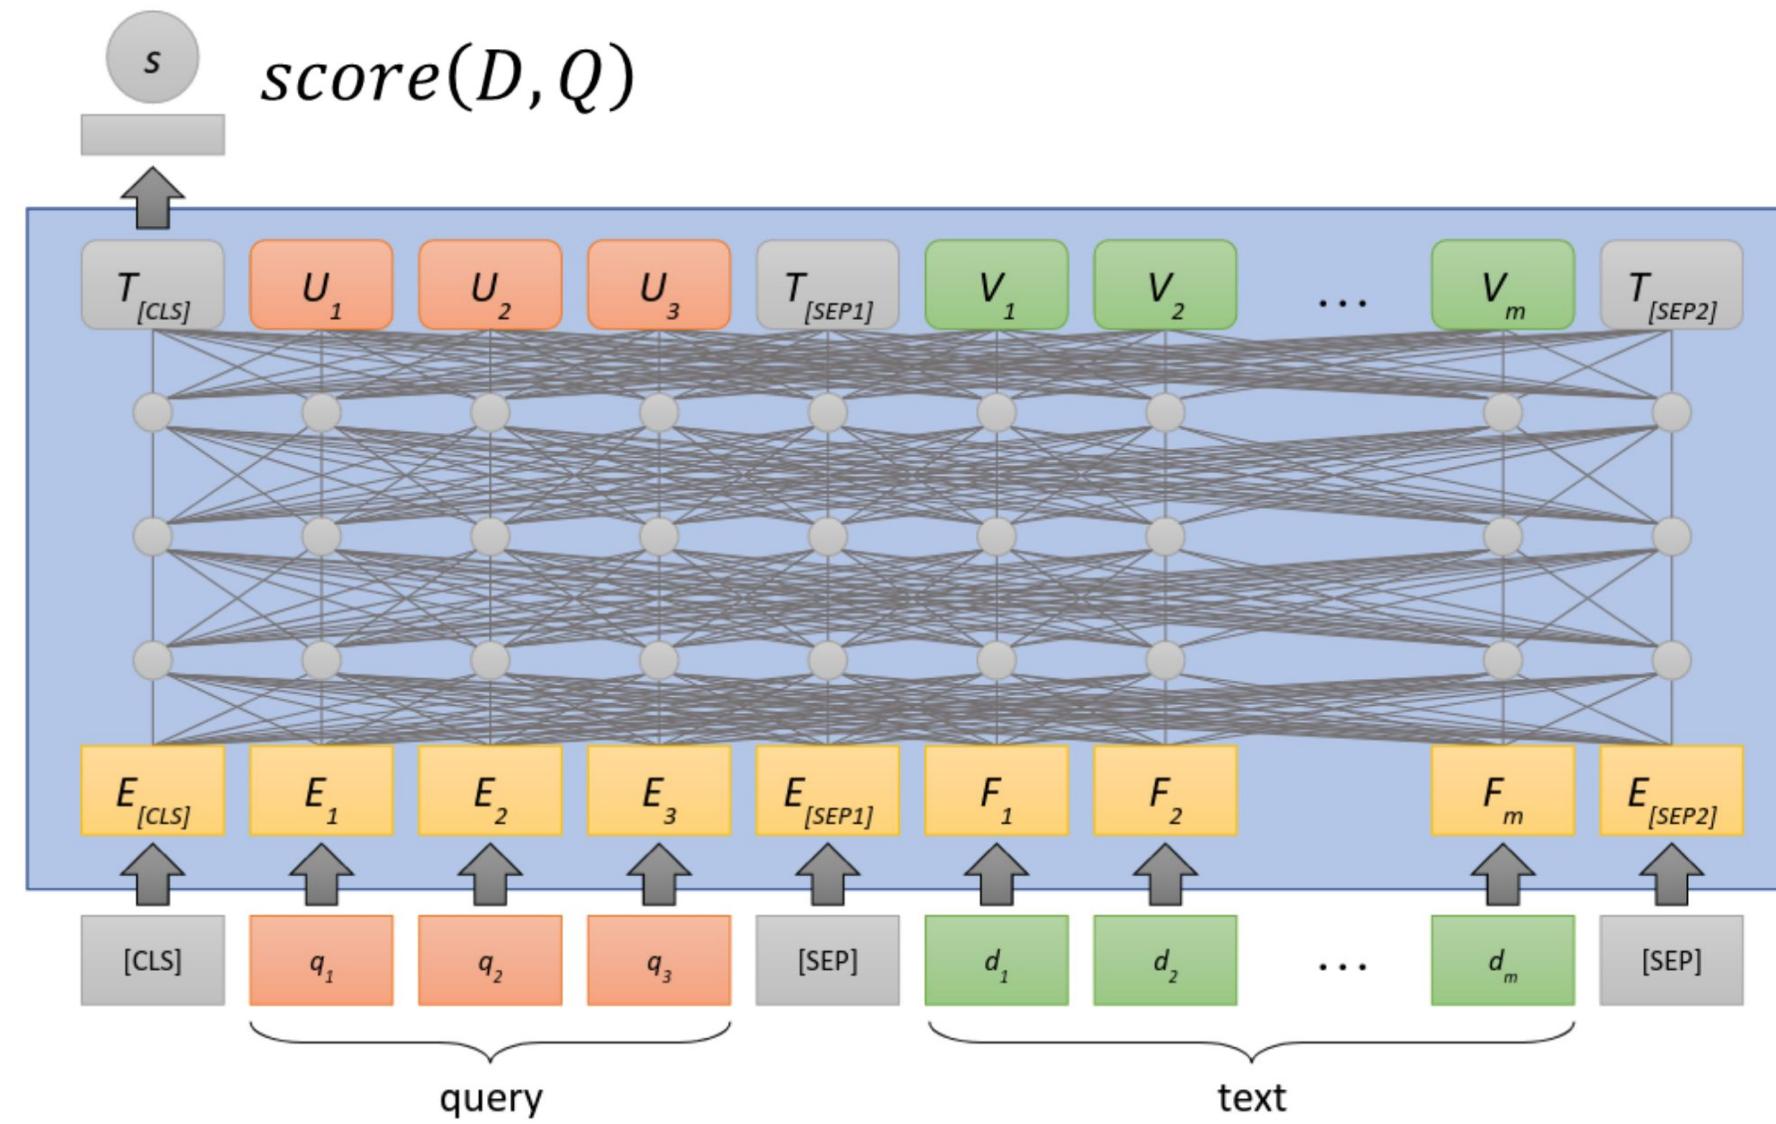
Cloud TPU v3 Pod
100+ petaflops
32 TB HBM

Lots of compute

Outline

- Introduction
- Ranking with soft matching
- Transformers & contextualized embeddings
- **Transformer approaches for ranking**
 - Re-ranking with cross-encoders
 - Learned sparse retrieval
 - Document expansion
- Conclusion

Three families of approaches

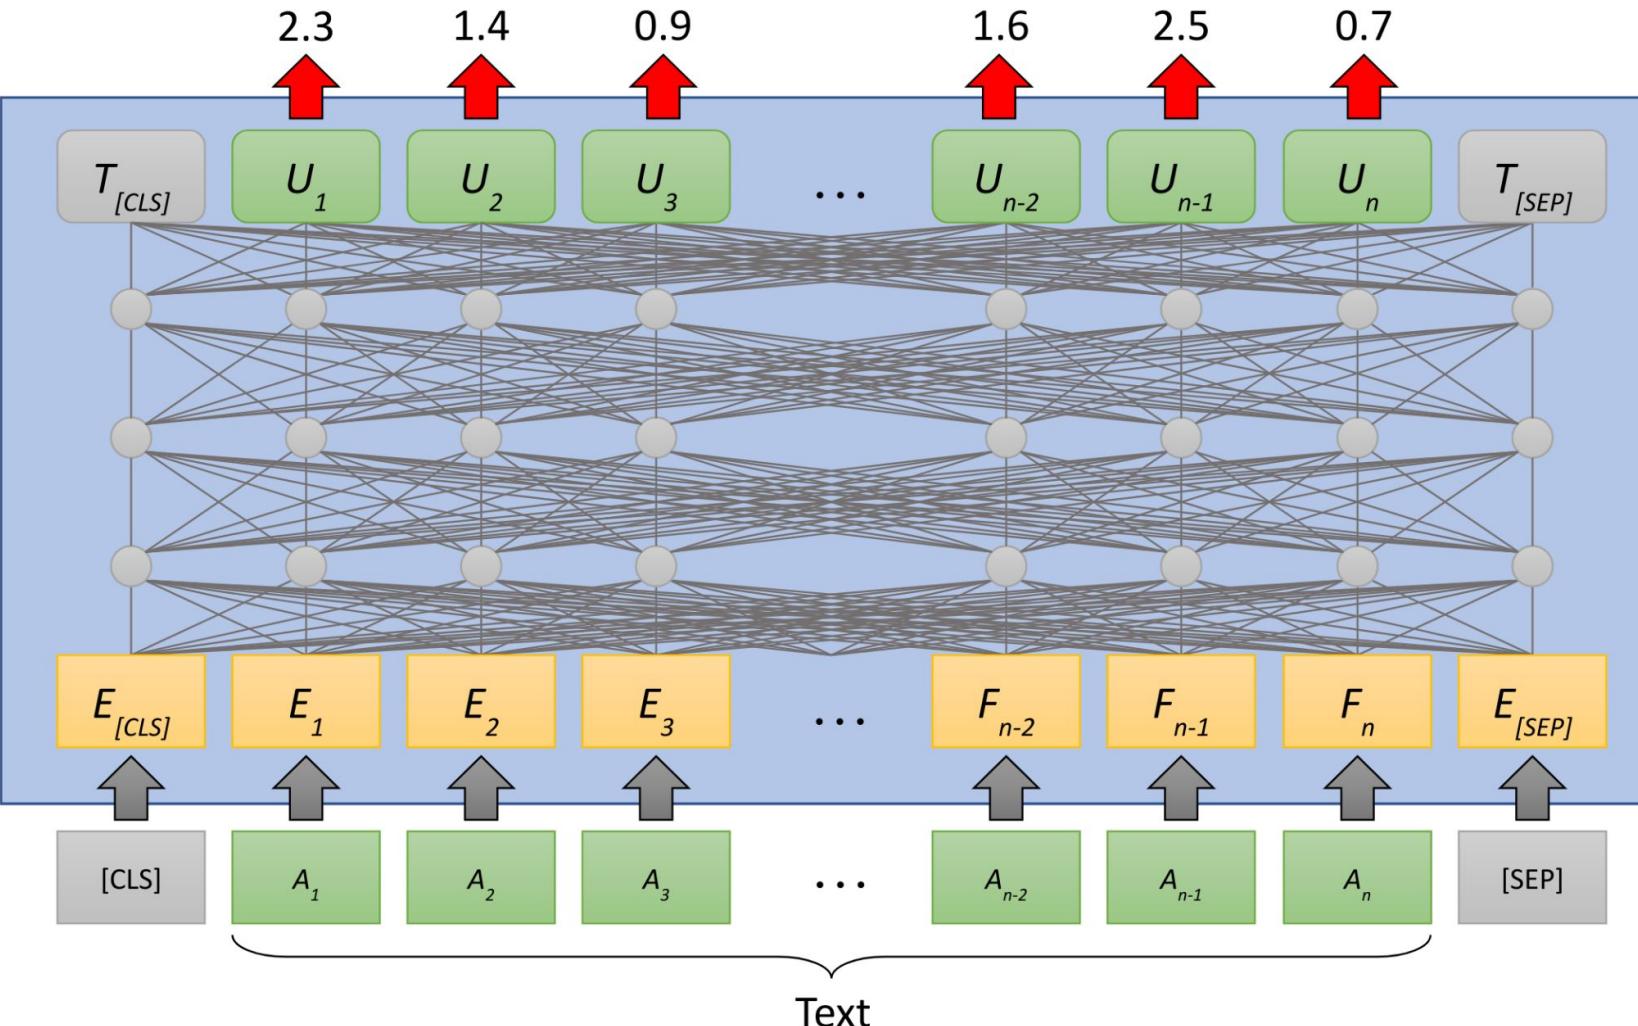


Cross-encoder

- Input: Q-D pair
 - Model outputs score
 - Slow but robust
- (Family in previous results)

Data structure:
Forward index

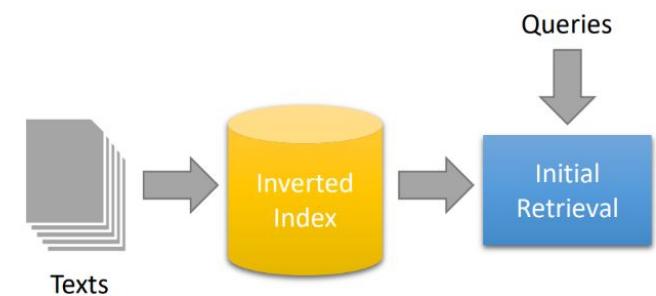
Three families of approaches



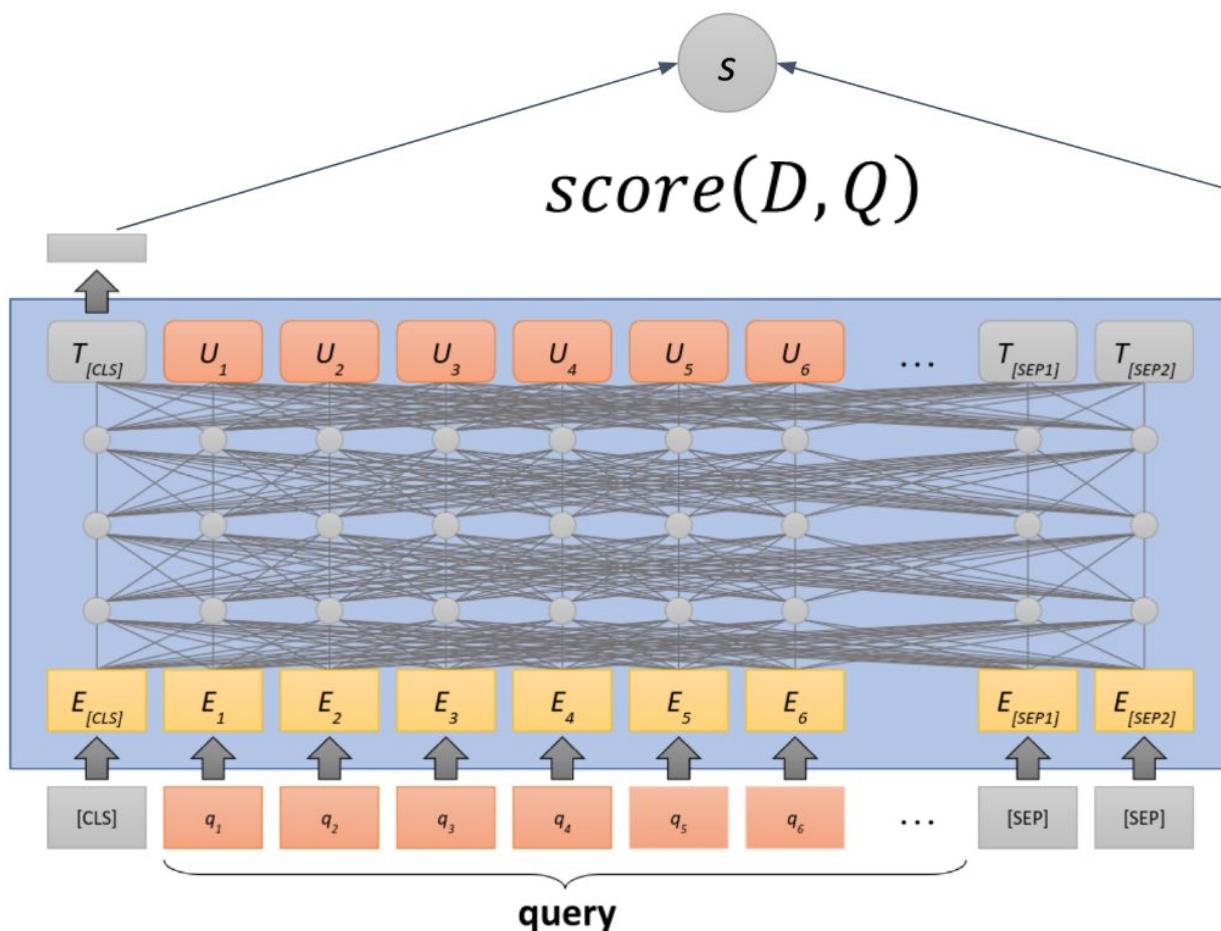
Learned sparse retrieval

- Input: Q or D
- Model outputs new term weights (replacing TF-IDF)
- Newest / least studied family

Data structure: Inverted index



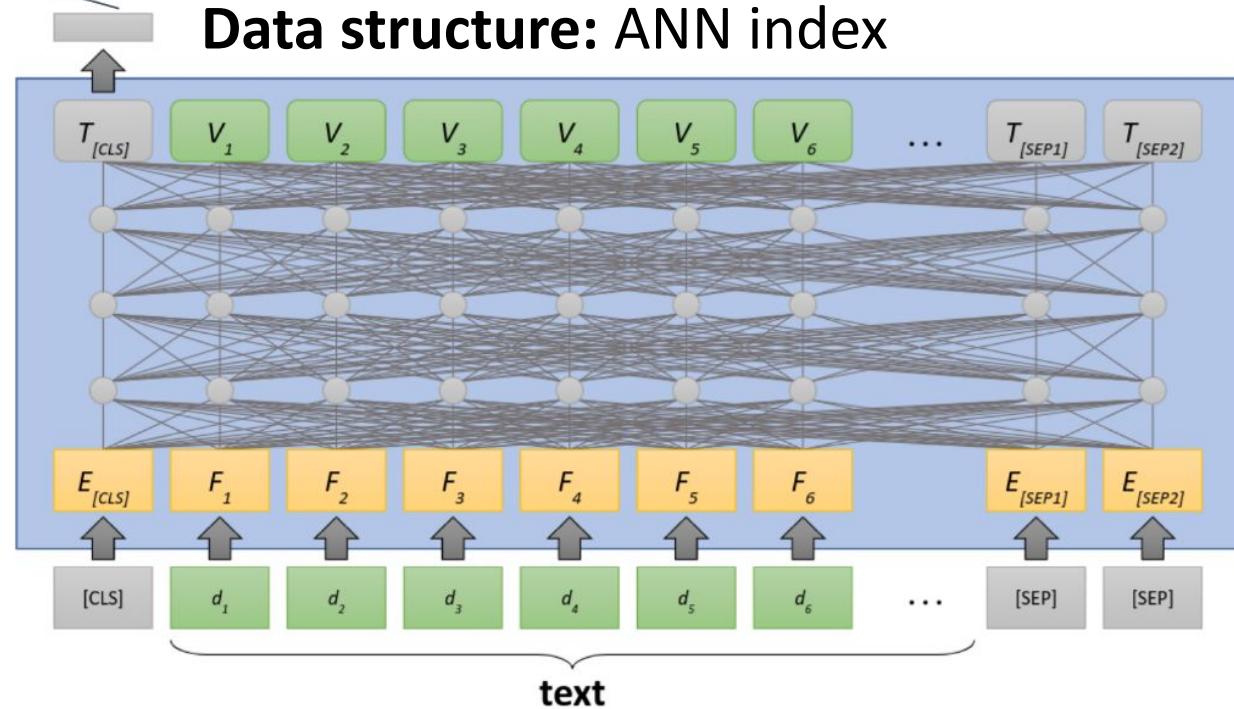
Three families of approaches



Bi-encoder

- Input: Q or D
- Model outputs vector
- Score by comparing Q, D vectors
- Faster, less effective & less robust

Data structure: ANN index



Outline

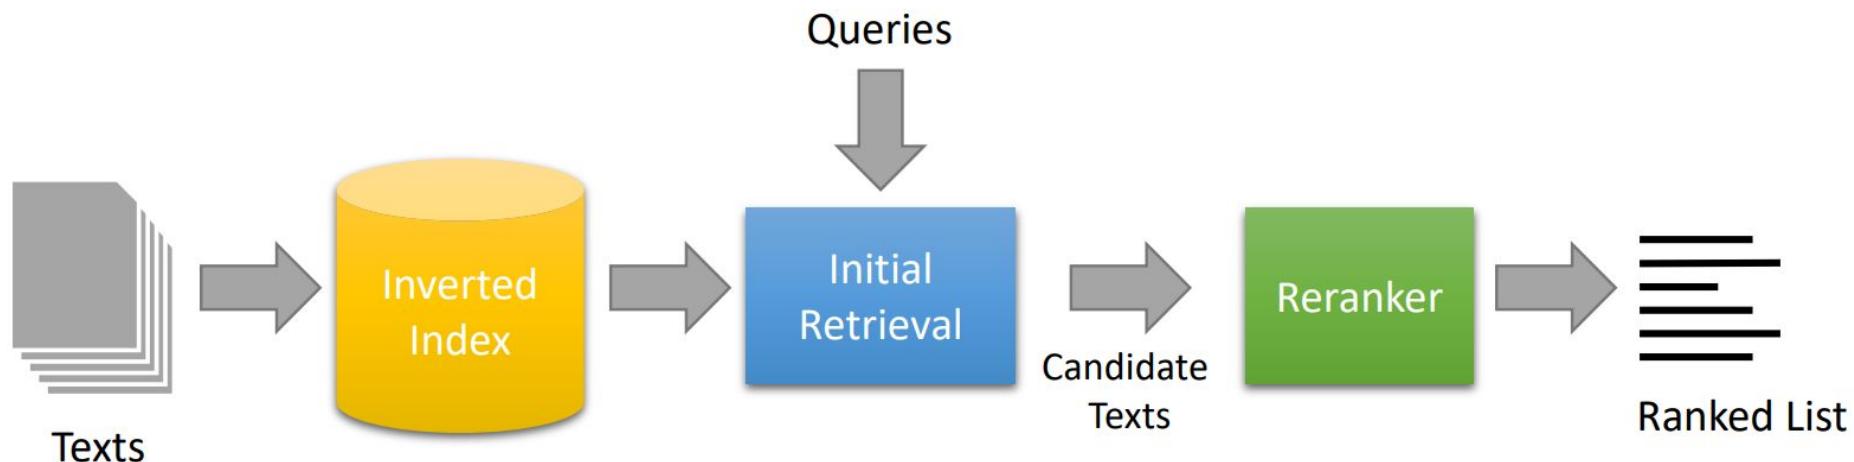
- Introduction
- Ranking with soft matching
- Transformers & contextualized embeddings
- Transformer approaches for ranking
 - **Re-ranking with cross-encoders**
 - Learned sparse retrieval
 - Document expansion
- Conclusion

Re-ranking with cross-encoders

How can we leverage a transformer's improved representations?

Transformer receives Q and D as input, then...

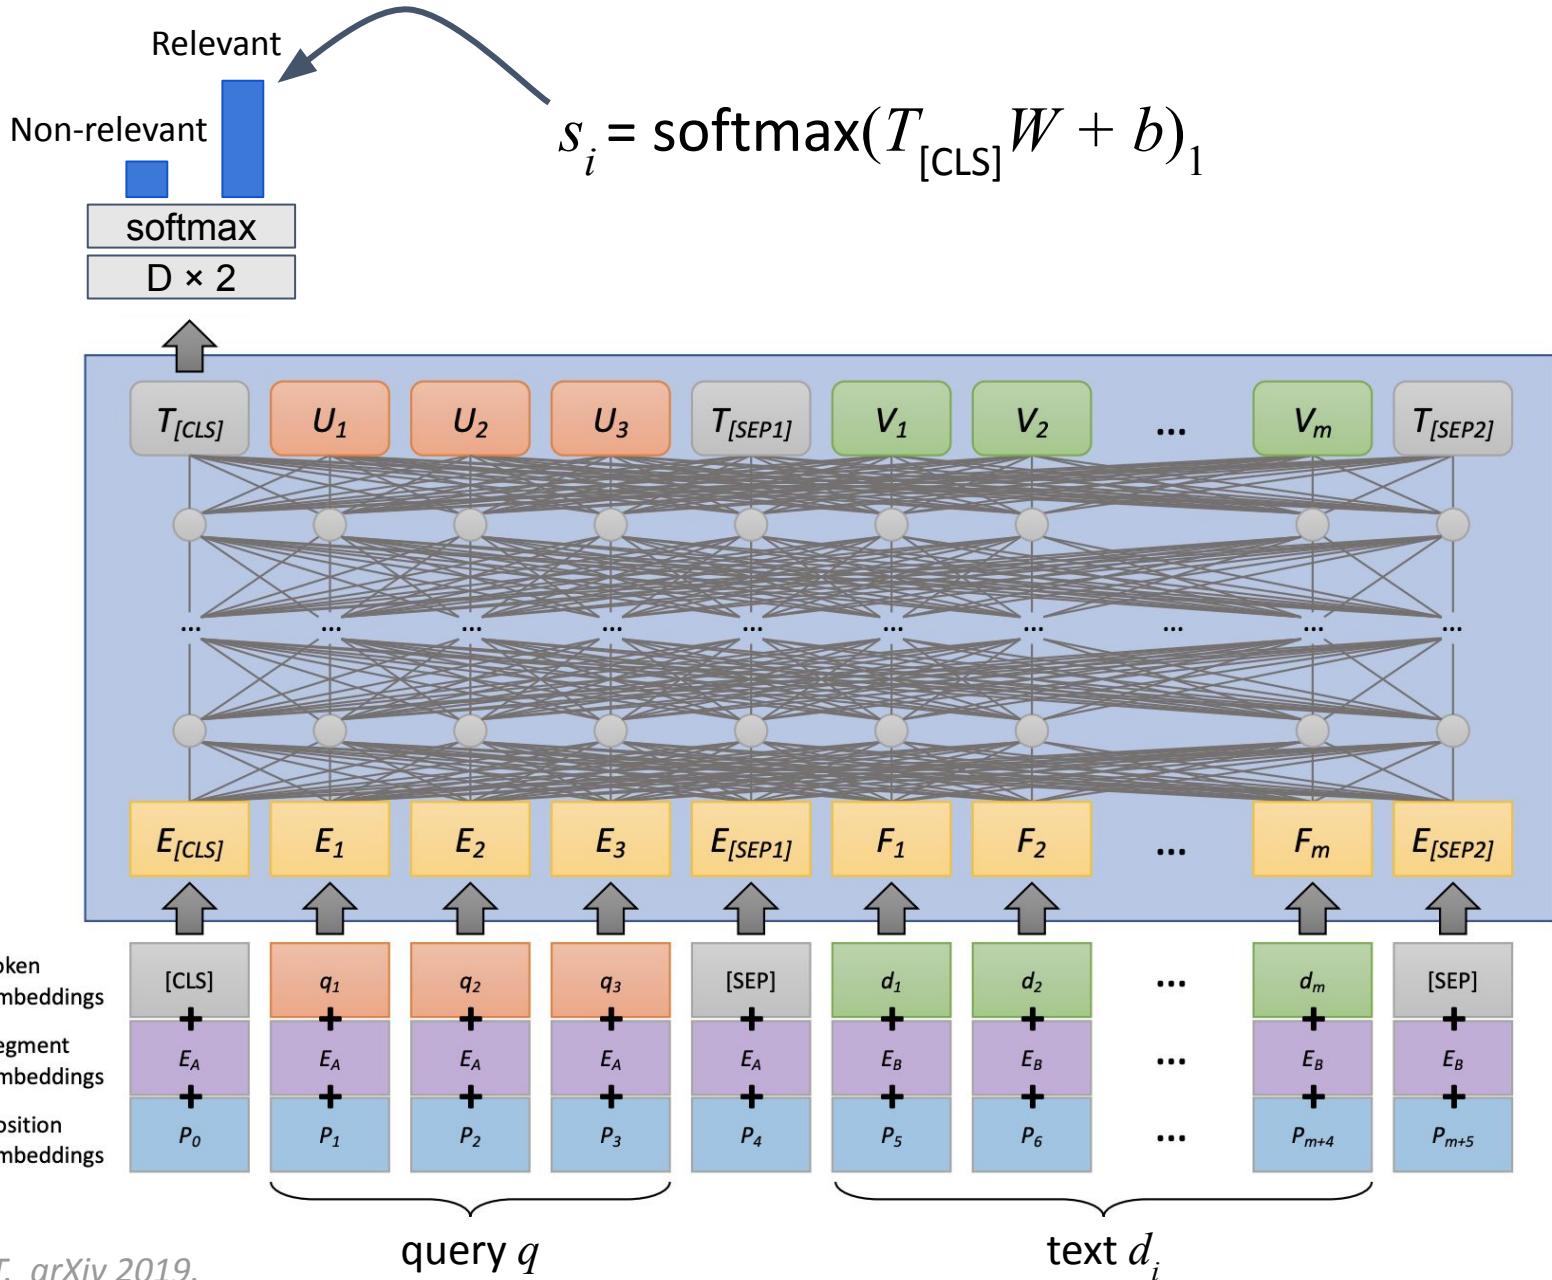
- monoBERT: predicts relevance score directly
- CEDR: produces contextualized embeddings



monoBERT reranker

We want:

$$s_i = P(\text{Relevant} = 1 | q, d_i)$$

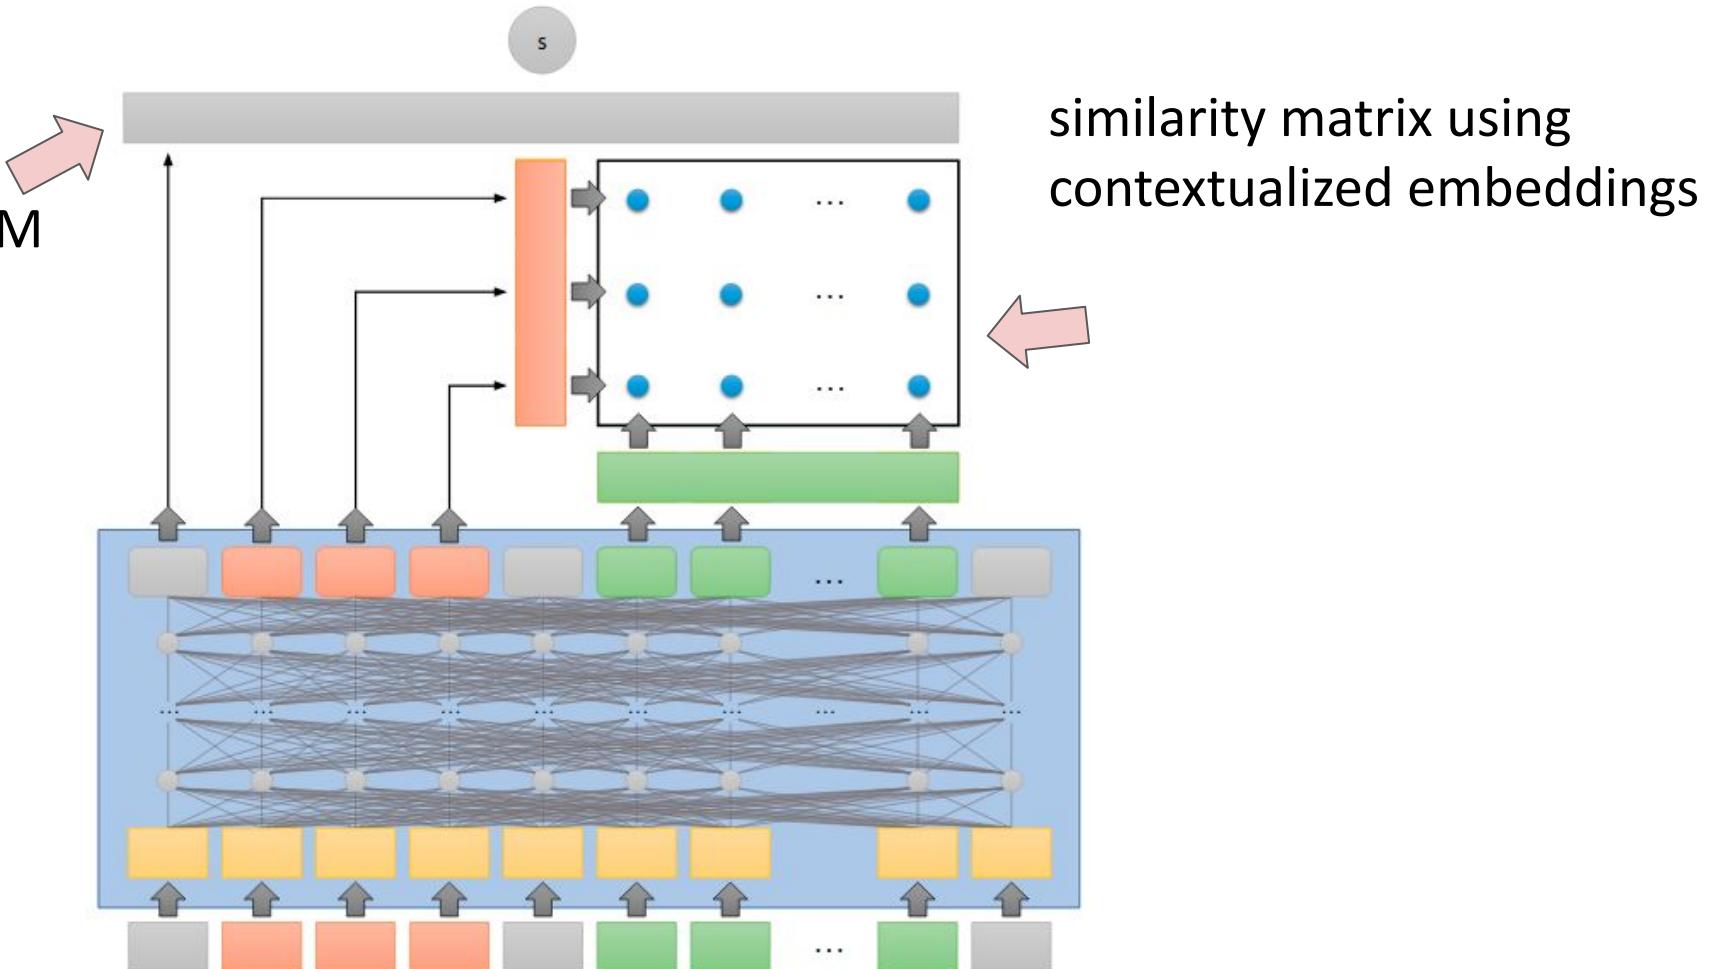


Results: monoBERT on TREC Deep Learning

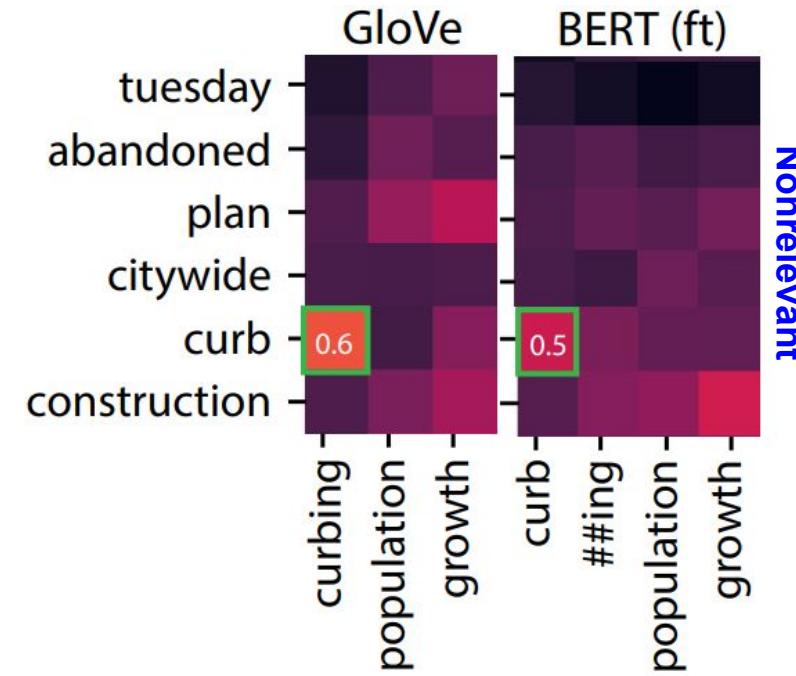
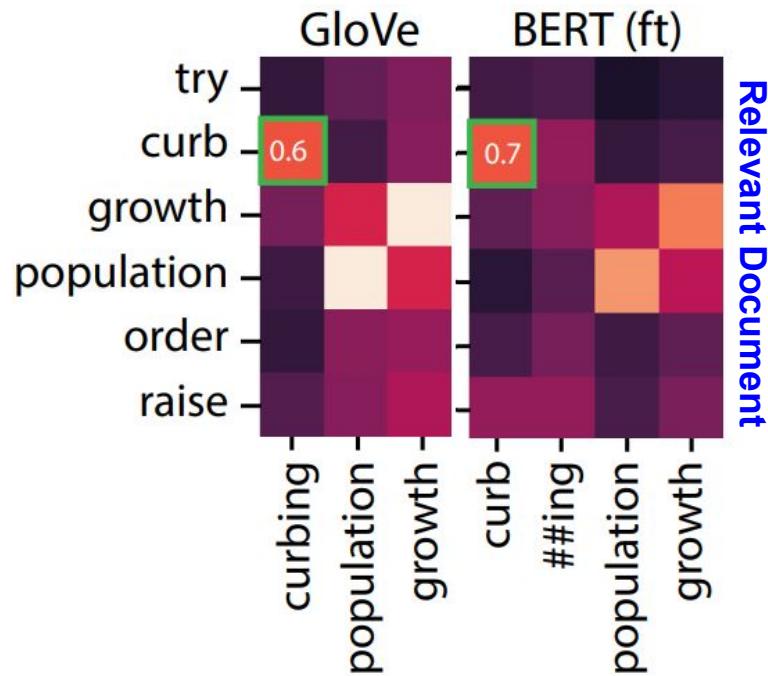
	nDCG@10	MAP	Recall@1k
BM25	0.506	0.377	0.739
+ monoBERT	0.738	0.506	0.739
BM25 + RM3	0.518	0.427	0.788
+ monoBERT	0.742	0.529	0.788

CEDR: Reranking with contextualized embeddings

pre-BERT model like DRMM



CEDR: Reranking with contextualized embeddings



Results: CEDR on Robust04

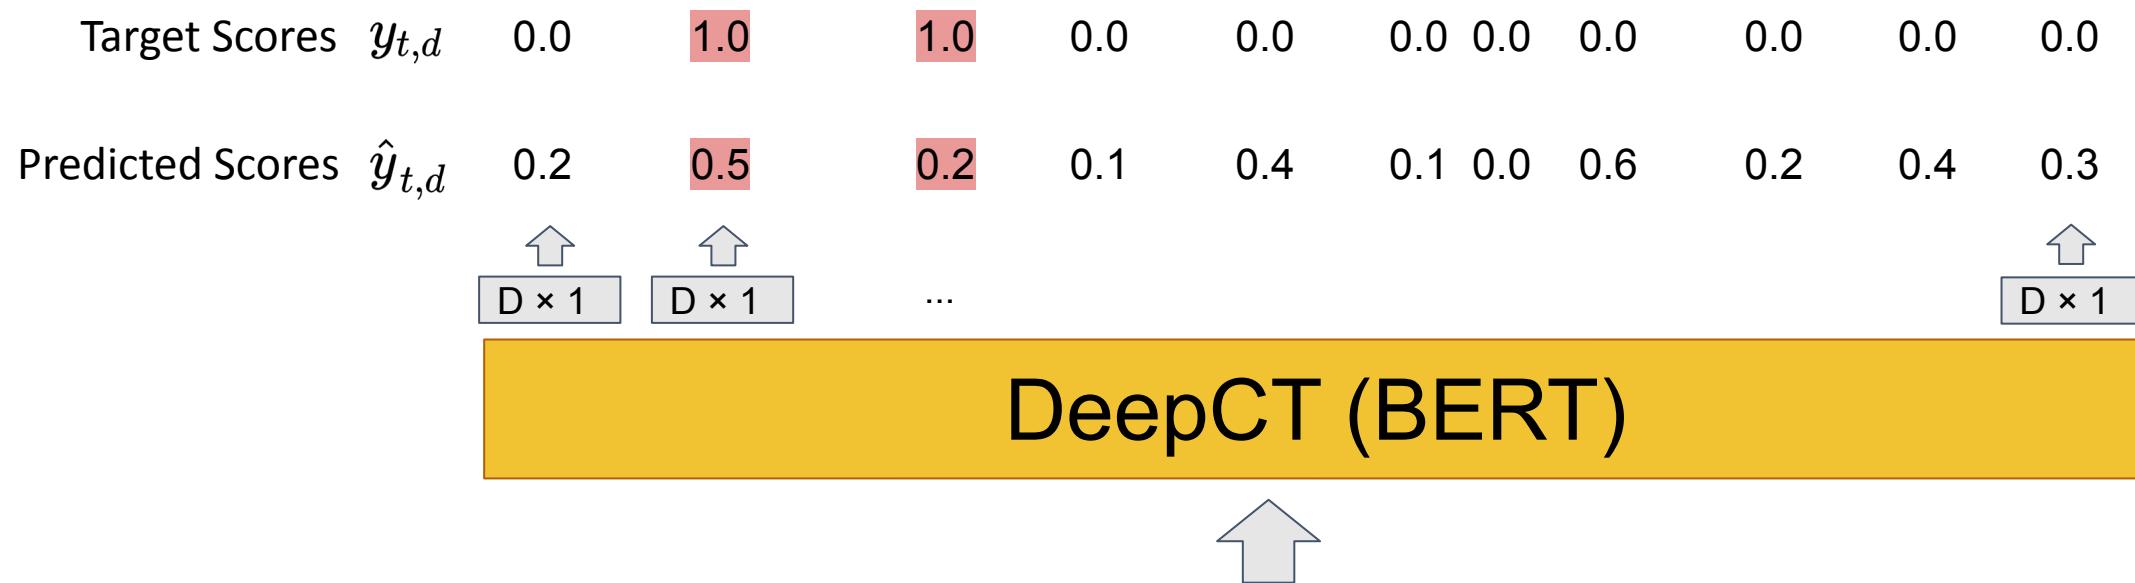
	Method	Input Representation	Robust04 nDCG@20
(1)	BM25	n/a	0.4140
(2)	Vanilla BERT	BERT (fine-tuned)	[B] 0.4541
(3a)	PACRR	GloVe	0.4043
(3b)	PACRR	BERT	0.4200
(3c)	PACRR	BERT (fine-tuned)	[BVG] 0.5135
(3d)	CEDR–PACRR	BERT (fine-tuned)	[BVG] 0.5150
(4a)	KNRM	GloVe	0.3871
(4b)	KNRM	BERT	[G] 0.4318
(4c)	KNRM	BERT (fine-tuned)	[BVG] 0.4858
(4d)	CEDR–KNRM	BERT (fine-tuned)	[BVGN] 0.5381
(5a)	DRMM	GloVe	0.3040
(5b)	DRMM	BERT	0.3194
(5c)	DRMM	BERT (fine-tuned)	[G] 0.4135
(5d)	CEDR–DRMM	BERT (fine-tuned)	[BVGN] 0.5259

Outline

- Introduction
- Ranking with soft matching
- Transformers & contextualized embeddings
- Transformer approaches for ranking
 - Re-ranking with cross-encoders
 - **Learned sparse retrieval**
 - Document expansion
- Conclusion

DeepCT

$$\text{loss} = \sum_t (\hat{y}_{t,d} - y_{t,d})^2$$



Text d : The Geocentric Theory was proposed by the greeks under the guidance...

Relevant query q : "who proposed the geocentric theory"

Indexing with DeepCT

Index



New document: "Researchers Researchers ... finding finding ... that that ... cinnamon cinnamon ... reduces ... "

Term Frequencies: 10 0 20 10 90 80 90 100 40 20 0 10 40



Predicted Scores $\hat{y}_{t,d}$ 0.1 0.0 0.2 0.1 0.9 0.8 0.9 1.0 0.4 0.2 0.0 0.1 0.4

$D \times 1$ $D \times 1$...

$D \times 1$

DeepCT (BERT)



Text d : Researchers are finding that cinnamon reduces blood sugar levels naturally when taken daily...

Results: DeepCT on MS MARCO

Model	MRR@10	R@1000
BM25	0.184	0.853
+ monoBERT	0.372	0.853
DeepCT	0.243	0.913

SPLADE: Sparse Lexical and Expansion Model

Key improvement: leverage the **MLM head** for weighting and expansion

→ Recall that MLM head predicts what term should fill in a [MASK]

Input: a query Q or document D

Output: a sparse vector of dimension $|V|$, to be indexed

1. For each term in the input, use MLM head to predict scores for $|V|$ terms
2. For each term in the vocabulary V, take maximum score as the term's weight
3. Use weights from #2 to represent the input (Q or D) as a $|V|$ vector; index

Results: SPLADE on MS MARCO

Model	MRR@10	R@1000
BM25	0.184	0.853
+ monoBERT	0.372	0.853
SPLADE	0.369	0.979
SPLADE-Doc	0.322	0.946

Outline

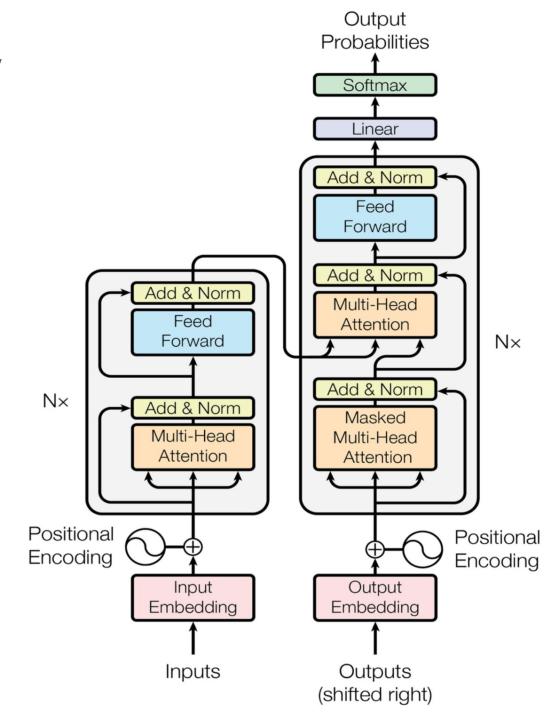
- Introduction
- Ranking with soft matching
- Transformers & contextualized embeddings
- Transformer approaches for ranking
 - Re-ranking with cross-encoders
 - Learned sparse retrieval
 - **Document expansion**
- Conclusion

Document expansion: doc2query

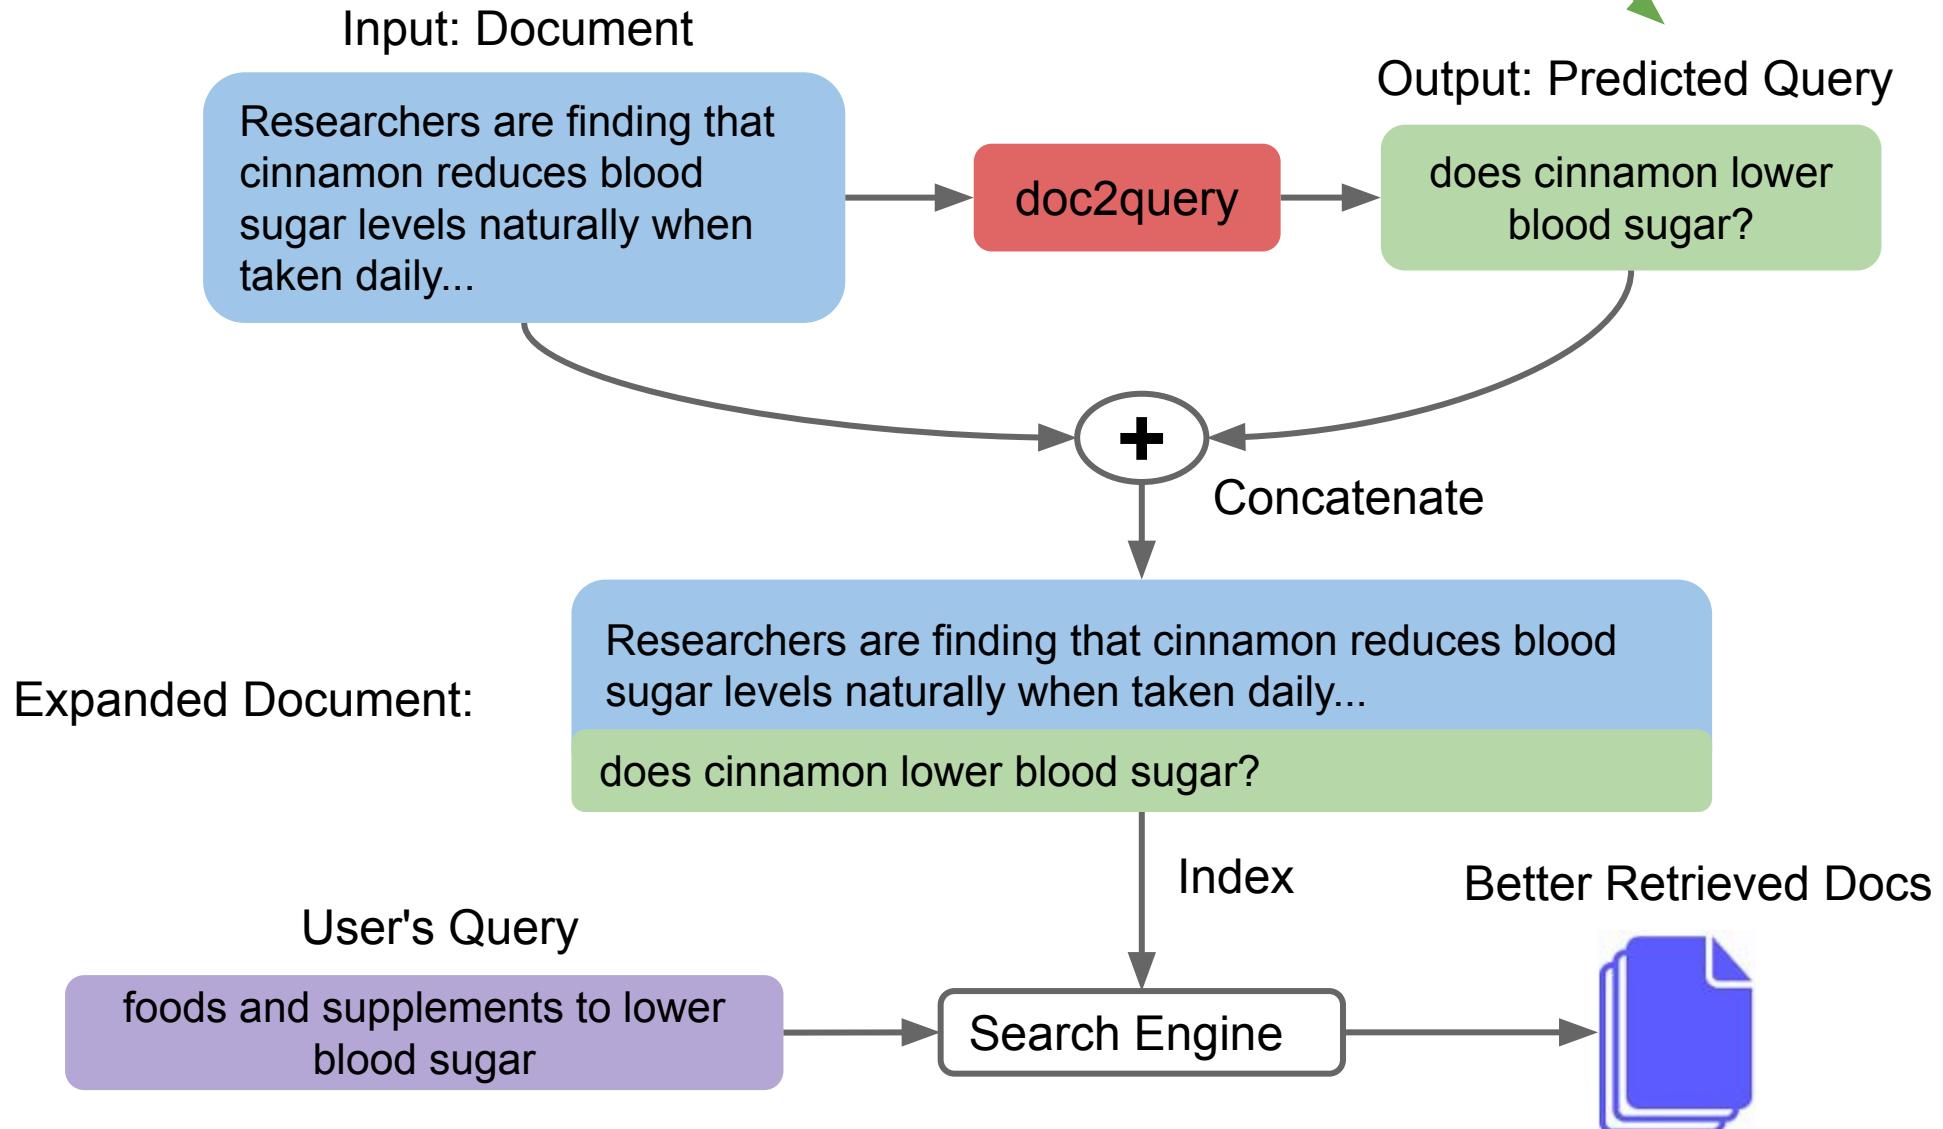
Idea: **generate** possible queries from a given document,
then use them to **expand** the document



Supervised training:
pairs of <query, relevant document>



doc2query



Examples from MS MARCO

Input Document: July is the hottest month in Washington DC with an average temperature of 27C (80F) and the coldest is January at 4C (38F) with the most daily sunshine hours at 9 in July. The wettest month is May with an average of 100mm of rain.

Predicted Query: weather in washington dc

Target query: what is the temperature in washington

Input Document: The Delaware River flows through Philadelphia into the Delaware Bay. It flows through an aqueduct in the Roundout Reservoir and then flows through Philadelphia and New Jersey before emptying into the Delaware Bay.

Predicted Query: what river flows through delaware

Target Query: where does the delaware river start and end

Input Document: sex chromosome - (genetics) a chromosome that determines the sex of an individual; mammals normally have two sex chromosomes
chromosome - a threadlike strand of DNA in the cell nucleus that carries the genes in a linear order; humans have 22 chromosome pairs plus two sex chromosomes.

Predicted Query: what is the relationship between genes and chromosomes

Target Query: which chromosome controls sex characteristics

Excluding stop-words:

69% copied

31% new

Results: doc2query on MS MARCO

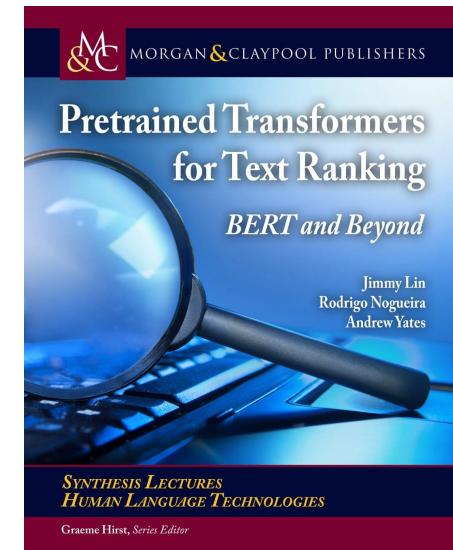
	MRR@10	R@1000
Original Document	.184	.853
+ Expansion New Words	.195	.907
+ Expansion Copied Words	.221	.893
+ Expansion Copied + New	.277	.944
Only Expansions (no original document)	.263	.927

Outline

- Introduction
- Ranking with soft matching
- Transformers & contextualized embeddings
- Transformer approaches for ranking
 - Re-ranking with cross-encoders
 - Learned sparse retrieval
 - Document expansion
- **Conclusion**

Conclusion

- Neural IR approaches can substantially improve search
- Learned sparse retrieval approaches compatible with an inverted index
- Cross-encoder approaches for reranking
- Dense retrieval (bi-encoders) next

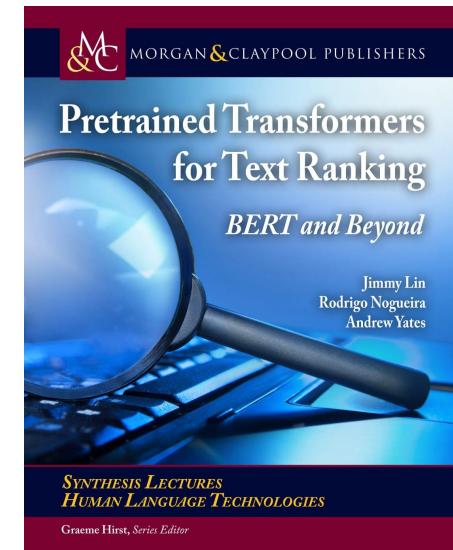


<https://arxiv.org/abs/2010.06467>
<https://bit.ly/tr4tr-book>

Conclusion

- Neural IR approaches can substantially improve search
- Learned sparse retrieval approaches compatible with an inverted index
- Cross-encoder approaches for reranking
- Dense retrieval (bi-encoders) next

Thanks!



<https://arxiv.org/abs/2010.06467>
<https://bit.ly/tr4tr-book>

