

# Representation Learning

Nils Reimers  
Principal Scientist @ cohere.ai  
[www.nils-reimers.de](http://www.nils-reimers.de)

co:here

# My Career Path



UBIQUITOUS  
KNOWLEDGE  
PROCESSING

Ph.D. + PostDoc



Neural Search Science Team  
Team Lead



sentence-transformers



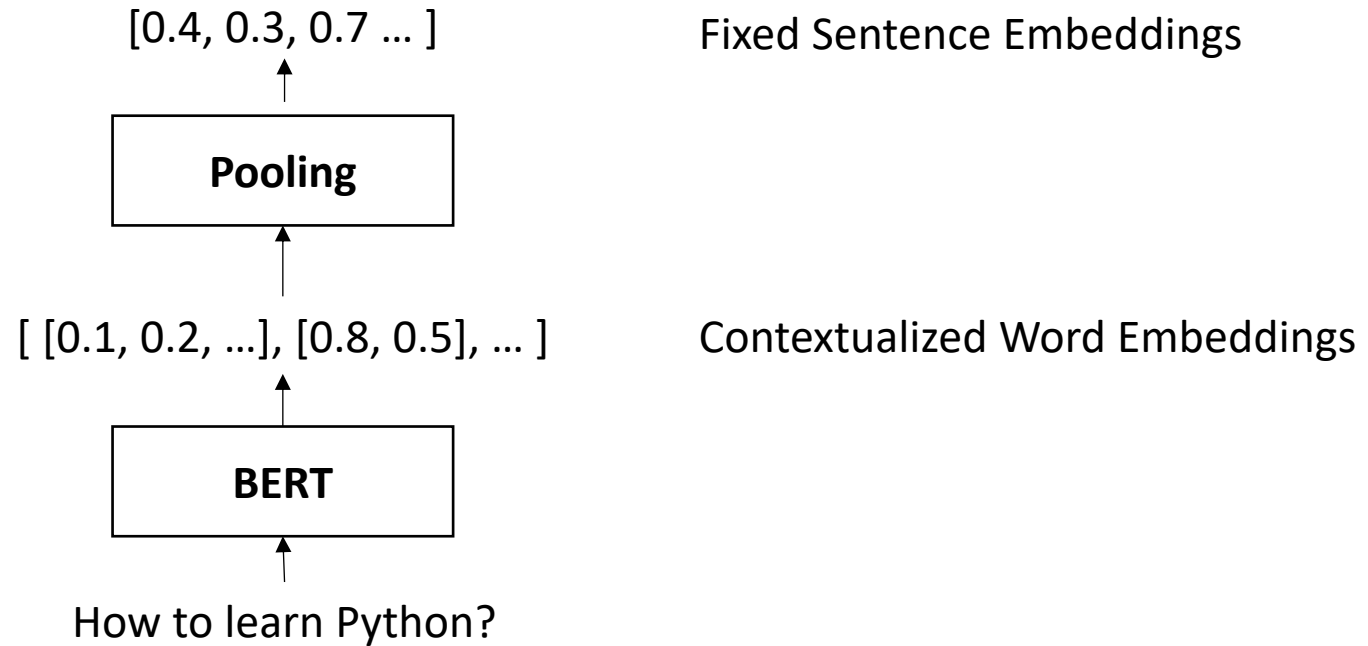
TSDAE

GPL


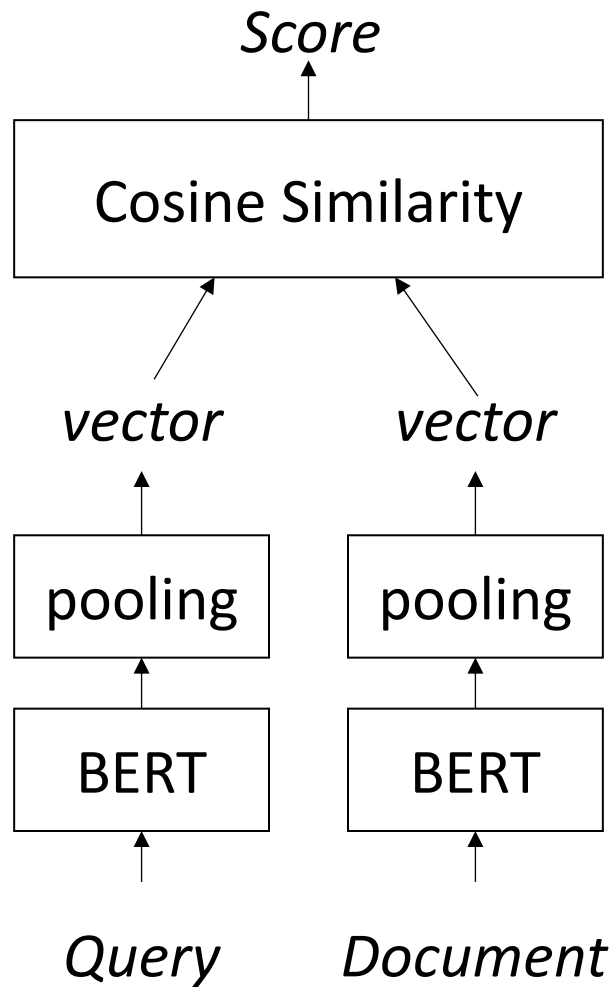
co:here

Principal Scientist / Director of Machine Learning  
Using very Large Language Models for search

# Sentence Embeddings Model




# Neural Search – Bi-Encoders

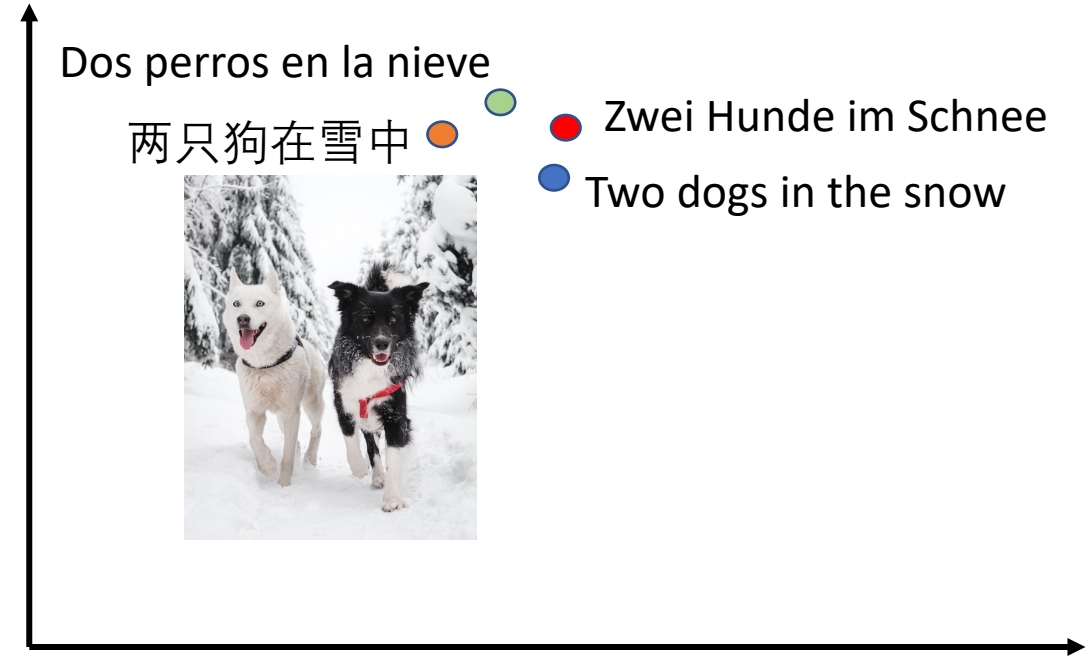
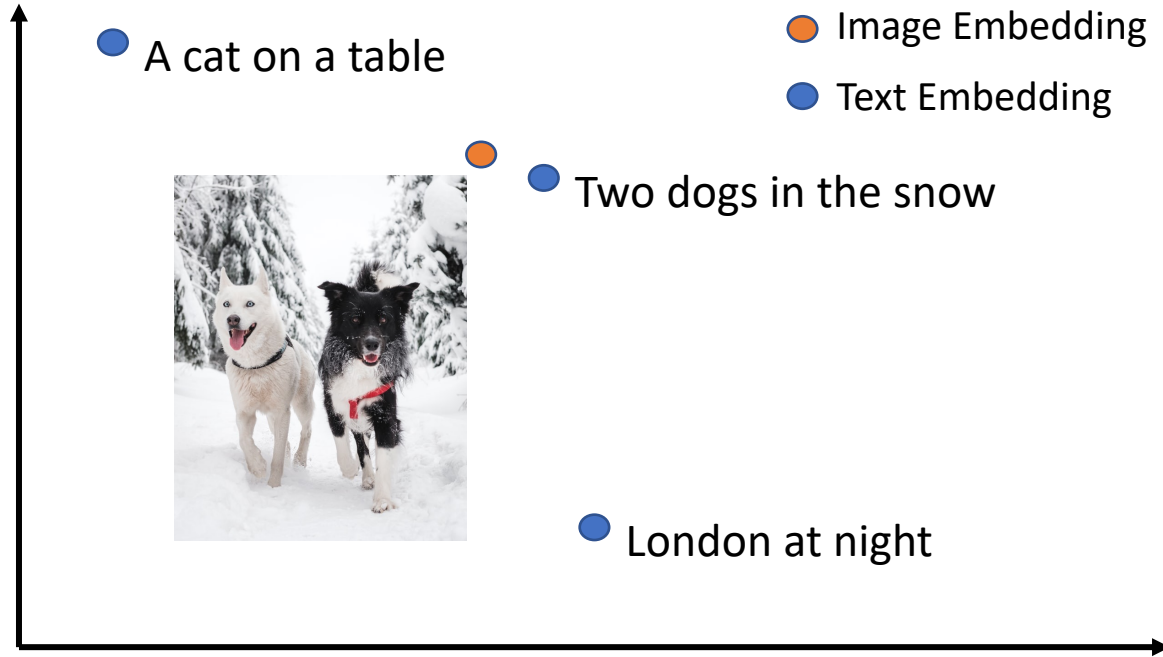


# Neural Search – Bi-Encoders

- Can overcome the lexical gap
  - US vs USA vs United States
- Respects the word order
  - Visa from Germany to Canada
  - Visa from Canada to Germany
- Knows about related terms
  - “spearman correlation numpy”  
finds the entry:  
“spearman correlation SciPy”




# Multi-Modal & Multi-Lingual Search



# Multiple Negative Ranking Loss

- Have positive pairs:
  - $(a_1, p_1)$
  - $(a_2, p_2)$
  - $(a_3, p_3)$
- Examples:
  - (query, answer-passage)
  - (question, duplicate\_question)
  - (paper title, cited paper title)
- $(a_i, p_i)$  should be close in vector space
- $(a_i, p_j)$  should be distant in vector space ( $i \neq j$ )
  - Unlikely that e.g. two randomly selected questions are similar
- Also called “training with in-batch negatives”, InfoNCE or NTXentLoss



# Multiple Negative Ranking Loss

- Mathematical Definition

$$L = -\frac{1}{n} \sum_{i=1}^n \frac{\exp(\text{sim}(a_i, p_i))}{\sum_j \exp(\text{sim}(a_i, p_j))}$$

- Sim: Similarity function between (a, p)
  - Cosine-Similarity
  - Dot-Product



# Multiple Negative Ranking Loss

## Intuitive Explanation

- $a_1$ : How many people live in Berlin?
  - $p_1$ : Around 3.5 million people live in Berlin
  - $p_2$ : Washington DC is the capital of the US
  - $p_3$ : The 2021 Olympics are held in Japan
- Compute text embeddings & compute similarities:
  - $\text{sim}(a_1, p_1) = 0.5$
  - $\text{sim}(a_1, p_2) = 0.3$
  - $\text{sim}(a_1, p_3) = 0.1$
- See it as classification task and use Cross-Entropy Loss:
  - Prediction:  $[0.5, 0.3, 0.1]$
  - Gold:  $[1, 0, 0]$

# Multiple Negative Ranking Loss

## Intuitive Explanation

- ( $a_1$ : How many people live in Berlin?,  $p_1$ : Around 3.5 million people live in Berlin)  
( $a_2$ : What is the capital of the US?,  $p_2$ : Washington DC is the capital of the US)  
( $a_3$ : Where are the Olympics this year?,  $p_3$ : The 2021 Olympics are held in Japan)

- Compute text embeddings & compute similarities:

$$\text{sim}(\text{vec\_a}, \text{vec\_b}) = \text{vec\_a} * \text{vec\_b}^T = \begin{bmatrix} \text{sim}(a_1, p_1), \text{sim}(a_1, p_2), \text{sim}(a_1, p_3) \\ \text{sim}(a_2, p_1), \text{sim}(a_2, p_2), \text{sim}(a_2, p_3), \\ \text{sim}(a_3, p_1), \text{sim}(a_3, p_2), \text{sim}(a_3, p_3) \end{bmatrix}$$

- See it as classification task and use Cross-Entropy Loss:

- Gold:  $\begin{bmatrix} 1, & 0, & 0, \\ 0, & 1, & 0, \\ 0, & 0, & 1 \end{bmatrix}$

# Multiple Negatives Ranking Loss Code

```
scores = self.similarity_fct(embeddings_a, embeddings_b) * self.scale
labels = torch.tensor(range(len(scores)), dtype=torch.long, device=scores.device) # Example a[i] should match with b[i]
return self.cross_entropy_loss(scores, labels)
```

[https://github.com/UKPLab/sentence-transformers/blob/master/sentence\\_transformers/losses/MultipleNegativesRankingLoss.py](https://github.com/UKPLab/sentence-transformers/blob/master/sentence_transformers/losses/MultipleNegativesRankingLoss.py)

# Multiple Negatives Ranking Loss

## Similarity Functions

- How to compute  $\text{sim}(a, b)$ ?
  - $a, b$  are vectors
  - Dot-product:  $\text{dot\_prod}(a, b) = ab^T$
  - Cosine-Similarity:  $\text{cos\_sim}(a, b) = (ab^T) / (||a|| ||b||)$ 
    - Does not work well, scores differences are too small
  - Scaled Cosine-Similarity:  $\text{scaled\_cos\_sim}(a, b) = C * \text{cos\_sim}(a, b)$ 
    - Works well with e.g.  $C=20$
  - Scaled dot-product:  $\text{scaled\_dot\_prod}(a, b) = C * \text{dot\_prod}(a, b)$

# Cosine-Similarity vs. Dot-Product


## Cosine-Similarity

- Vector has highest similarity to itself
  - $\text{cos\_sim}(a, a) = 1$
- With normalized vectors, equal to dot\_product
  - With max vector length = 1
- With normalized vectors, proportional to Euclidian distance
  - Works with k-means clustering

## Dot-Product

- Other vectors can have higher dot-product
  - $\text{dot}(a, a) < \text{dot}(a, b)$
- Might be slower with certain approximate nearest neighbor methods
  - Max vector length not know
- Does not work with k-means clust.

# Cosine-Similarity vs. Dot-Product



- Semantic search: Given short query, find longer passage
- Cosine-Similarity: Prefers retrieval of short passages close to query
- Dot-Product: Prefers longer passages (longer passage = longer vector = higher dot product)



# Optimizing the Multiple-Negatives-Ranking-Loss

- Training with  $\text{scaled\_cos\_sim}(a, b) = C * \text{cos\_sim}(a, b)$ 
  - How to choose the scale  $C$ ?  $\leq$  unclear, common values 14-20
  - ConveRT paper: Start at 1, end at 23, increase over first 10k steps
  - CLIP paper:  $\text{scaled\_cos\_sim}(a, b) = \exp(C) * \text{cos\_sim}(a, b)$  with  $C$  a learnable parameter
  - Unclear impact
    - Will it make a difference?
    - Does it depend on the data / task?
- Symmetric Multiple-Negatives-Ranking-Loss
  - Used in CLIP Paper
  - Compute:  $(\text{Loss}(A, P) + \text{Loss}(P, A)) / 2$
  - Swap anchor & positives (e.g. given answer, what is the question?)
  - Unclear impact

# Multiple-Negatives-Ranking-Loss with Additive Margin

$$\blacksquare \text{sim}(a_i, p_j) = \begin{cases} \text{sim}(a_i, p_i) - m & \text{if } i = j \\ \text{sim}(a_i, p_j) & \end{cases}$$

- Subtract value  $m$  from positive pairs
  - Cosine-similarity with margin 0.3 used in LaBSE paper with translation pairs
- Unclear impact of margin for other tasks / datasets




- Used in: LaBSE: <https://arxiv.org/abs/2007.01852> & <https://arxiv.org/abs/1902.08564>



# Multiple Negative Ranking Loss

## Hard Negatives

- Larger batch size => task more difficult => better results
  - Given query, which of the 10 passages provide the answer?
  - Given query, which of the 1k passages provide the answer?



# Multiple Negative Ranking Loss


## Hard Negatives

- Train with tuples:

$(a_1, p_1, n_1)$

$(a_2, p_2, n_2)$

- $n_i$  should be similar to  $p_i$  but not match with  $a_i$
- Bad example:
  - a: How many people live in London?
  - p: Around 9 million people live in London
  - n: London has a population of 9 million people.
- Good example:
  - a: How many people live in London?
  - p: Around 9 million people live in London
  - n: Around 1 million people live in Birmingham, second to London.



# How to find hard-negatives?

- Quality of hard-negatives significantly improves the performance
- Finding good hard negatives not easy
- Strategy 1: Exploit structure in your data
  - Citation graph: (Title, Cited\_Paper, Paper\_Cited\_by\_Cited\_Paper)
  - Q&A: (Question, Answer with many stars, Answer with few stars)
- Strategy 2: Mine hard negative:
  - Use BM25 to find top-100 most similar texts to anchor / positive
  - Select one of these randomly
  - Make sure that these are actually negatives!

# Improving Quality with Better Batches

- Assume you have (question, answer) pairs from StackExchange
  - 140 different subforums: StackOverflow, Travel, Cooking, ...
- Naïve approach:
  - Randomly sample data from all pairs:  
[ (question\_python, answer\_python),  
 (question\_travel, answer\_travel),  
 (question\_pasta, answer\_pasta)]
  - Finding the right answer for a given question is easy
    - Question about Python => Take that one programming answer in the batch...

# Improving Quality with Better Batches

- Assume you have (question, answer) pairs from StackExchange
  - 140 different subforums: StackOverflow, Travel, Cooking, ...
- Better approach
  - Sample pairs from one subforum (e.g. StackOverflow)
    - [ (question\_python, answer\_python),  
    (question\_java,     answer\_java),  
    (question\_c,         answer\_c)]

# Improving Quality with Better Batches

- Assume you have (question, answer) pairs from StackExchange
  - 140 different subforums: StackOverflow, Travel, Cooking, ...
- Even better approach (?)
  - Sample pairs from same / similar tags (e.g. StackOverflow, Python tag)
    - [ (question\_python, answer\_python),  
 (question\_numpy, answer\_numpy),  
 (question\_pandas, answer\_pandas)]
- Adding random batches might still be needed
  - Otherwise StackOverflow vector space could overlap with Travel vector space
  - 90% difficult batches, 10% easy random batches
  - Or: start with mainly random batches, then go to difficult batches

# Bi-Encoders and the Curse of the Unknowns


- How do Bi-Encoders handle unknown words?
  - Not seen during pre-training
  - Not seen during fine-tuning
- Where to put these words in a vector space?
  - XLNet
  - Clexchain
  - Forwrd
  - 0xc004f213
- How to know
  - Corona Virus  $\Leftrightarrow$  COVID-19  $\Leftrightarrow$  SARS-Cov-2
  - Q: “Which vision transformer model is the best?”  
A: “ViT has been doing great in our experiments”

# Challenge of Unknown Words for Dense Bi-Encoders

“BigBirdPegasus”

Dense  
Encoder


- 1) bigboss
- 2) bigdata
- 3) bigass





# Unknown Words for Sparse Bi-Encoders

“BigBirdPegasus”



Split query in word pieces:

- big:2.1, ##bird: 2.0, ##pe: 1.8, ##gas: 2.0, ##us:1.9

Some related terms are added:

- ##birds: 1.2, giant: 0.7



# BEIR – Benchmarking IR

9 Tasks

18 Datasets



**Beir**  
Benchmarking IR

## Fact Checking

### FEVER



Wiki

QUERY Natural Claim  
DOCS Wikipedia Articles



Wiki

Climate-FEVER  
QUERY Climate-based Claim  
DOCS Wikipedia Articles



Scientific

SciFact  
QUERY Scientific claim  
DOCS PubMed Articles

## Citation-Prediction



Scientific

SCIDOCs  
QUERY Article Title  
DOCS PubMed Articles

## Dup. Question Retrieval



Quora

Quora  
QUERY Query Title  
DOCS Quora Questions



StackEx.

CQADupStack  
QUERY Query Title  
DOCS Query Title + Body

## Argument Retrieval



Misc.

Touche-2020  
QUERY Controversial Query  
DOCS Args.me Arguments



Misc.

ArguAna  
QUERY Argument  
DOCS Idebate Arguments

## News Retrieval



News

TREC-NEWS  
QUERY News Headline  
DOCS News Articles



News

Robust04  
QUERY News Query  
DOCS News Articles

## Question-Answering



Wiki

NQ  
QUERY Natural Query  
DOCS Wikipedia Articles



Wiki

HotpotQA  
QUERY Multi-Hop Query  
DOCS Wikipedia Articles



Finance

FiQA-2018  
QUERY Financial Query  
DOCS Investment Articles



Twitter

## Tweet Retrieval


Signal-1M  
QUERY News Headline  
DOCS Twitter Tweets

## Bio-Medical IR



Scientific

TREC-COVID  
QUERY COVID-19 Query  
DOCS CORD-19 Articles



Scientific

BioASQ  
QUERY Bio-Medical Query  
DOCS PubMed Articles



Scientific

NFCorpus  
QUERY Nutrition Facts  
DOCS PubMed Articles



Wiki

## Entity Retrieval


DBPedia  
QUERY Entity-based Query  
DOCS DBPedia Articles

# Bi-Encoders vs Lexical Search

Dataset	BM25	Dense Model (TAS-B)	Difference
In-Domain	22.8	40.8	+18.0
BioASQ	46.5	38.3	-8.2
SCIDOCS	15.9	14.9	-1.0


- BM25 was better on 10 / 18 datasets

# Do Models Generalize?



- BM25 lexical search a strong baseline
- BM25 + CrossEncoder re-ranking perform the best
- Dense embedding models (TAS-B, ANCE, DPR) with issues for unknown domains
- Sparse embedding models (SPLADEv2) better for unknown domains

# Cross-Encoders vs Bi-Encoders



# Cross-Encoders vs Bi-Encoders

<b>Dataset</b>	<b>BM25</b>	<b>Dense Model (TAS-B)</b>	<b>BM25 + CE</b>
In-Domain	22.8	40.8	41.3
BioASQ	46.5	38.3	52.3
CQADupStack	29.9	31.4	37.0
TREC-COVID	65.5	48.1	75.7
SCIDOCS	15.9	14.9	16.6


- BM25 + CE on average 13.8 points better than dense

# Why not using Cross-Encoders / doc2query?

- Cross-Encoders are slow (even small ones)
  - E.g. query has 10 tokens, docs have 240 tokens, re-rank 100 docs
  - Bi-Encoders: Compute embedding for query (e.g. 10ms)
  - Cross-Encoder: Re-rank 100 x 250 token docs
    - Forward pass for 250 tokens takes  $\sim 25 \times 25 = 625$  times longer
    - Overall 62,500 times longer to get results
- Doc2query is slow at indexing
  - Generates 40 query per passage
  - Question generation is extremely slow
  - Costs to generate queries for 8M docs: \$750
  - Computing dense embeddings: \$1

# How to Adapt Bi-Encoders to New Domains?


## Adaptive Pre-Training



✗ Requires (expensive) training on labeled source data

---


## What we want:



✓ Use pre-trained models and adapt to you domain

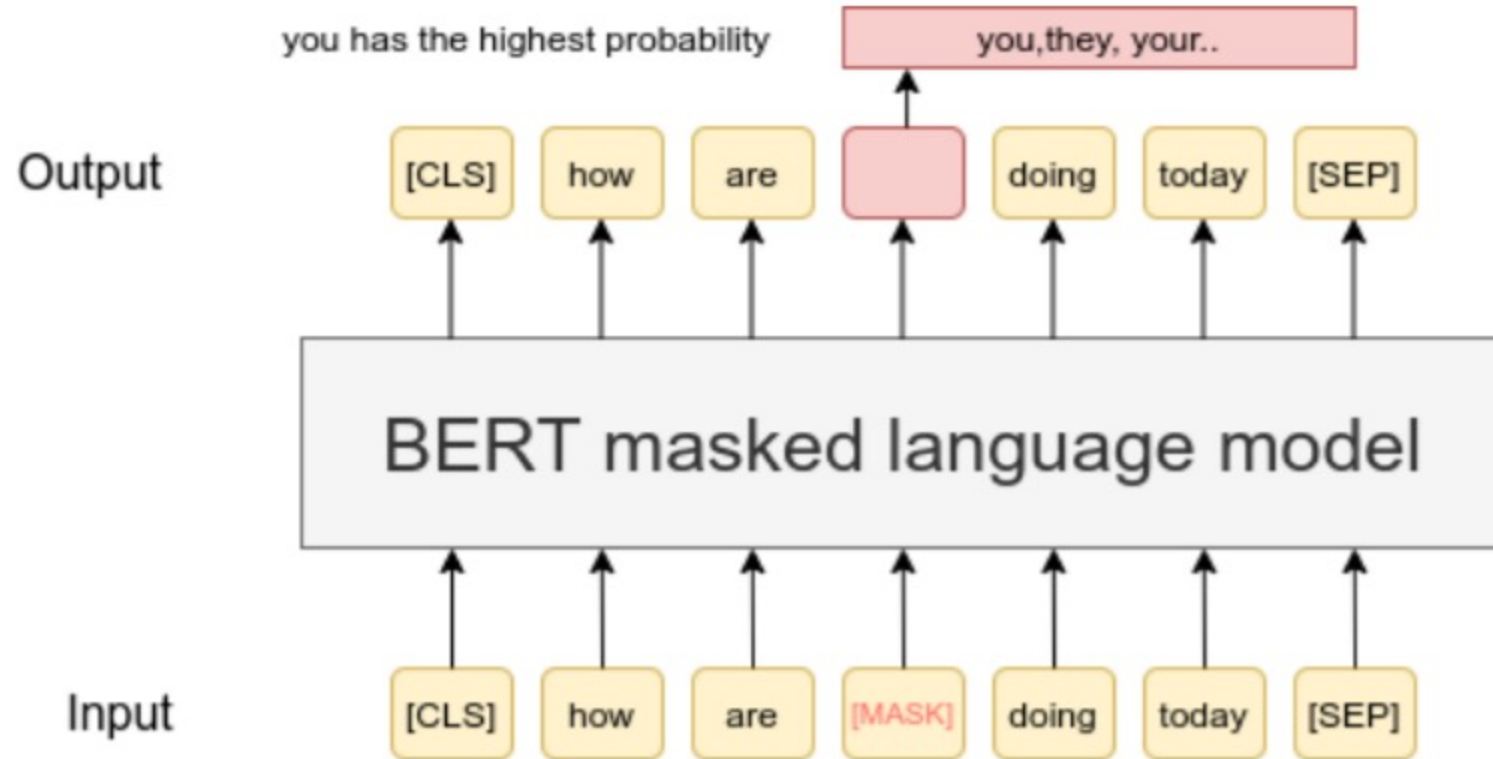


# Adaptive Pre-Training

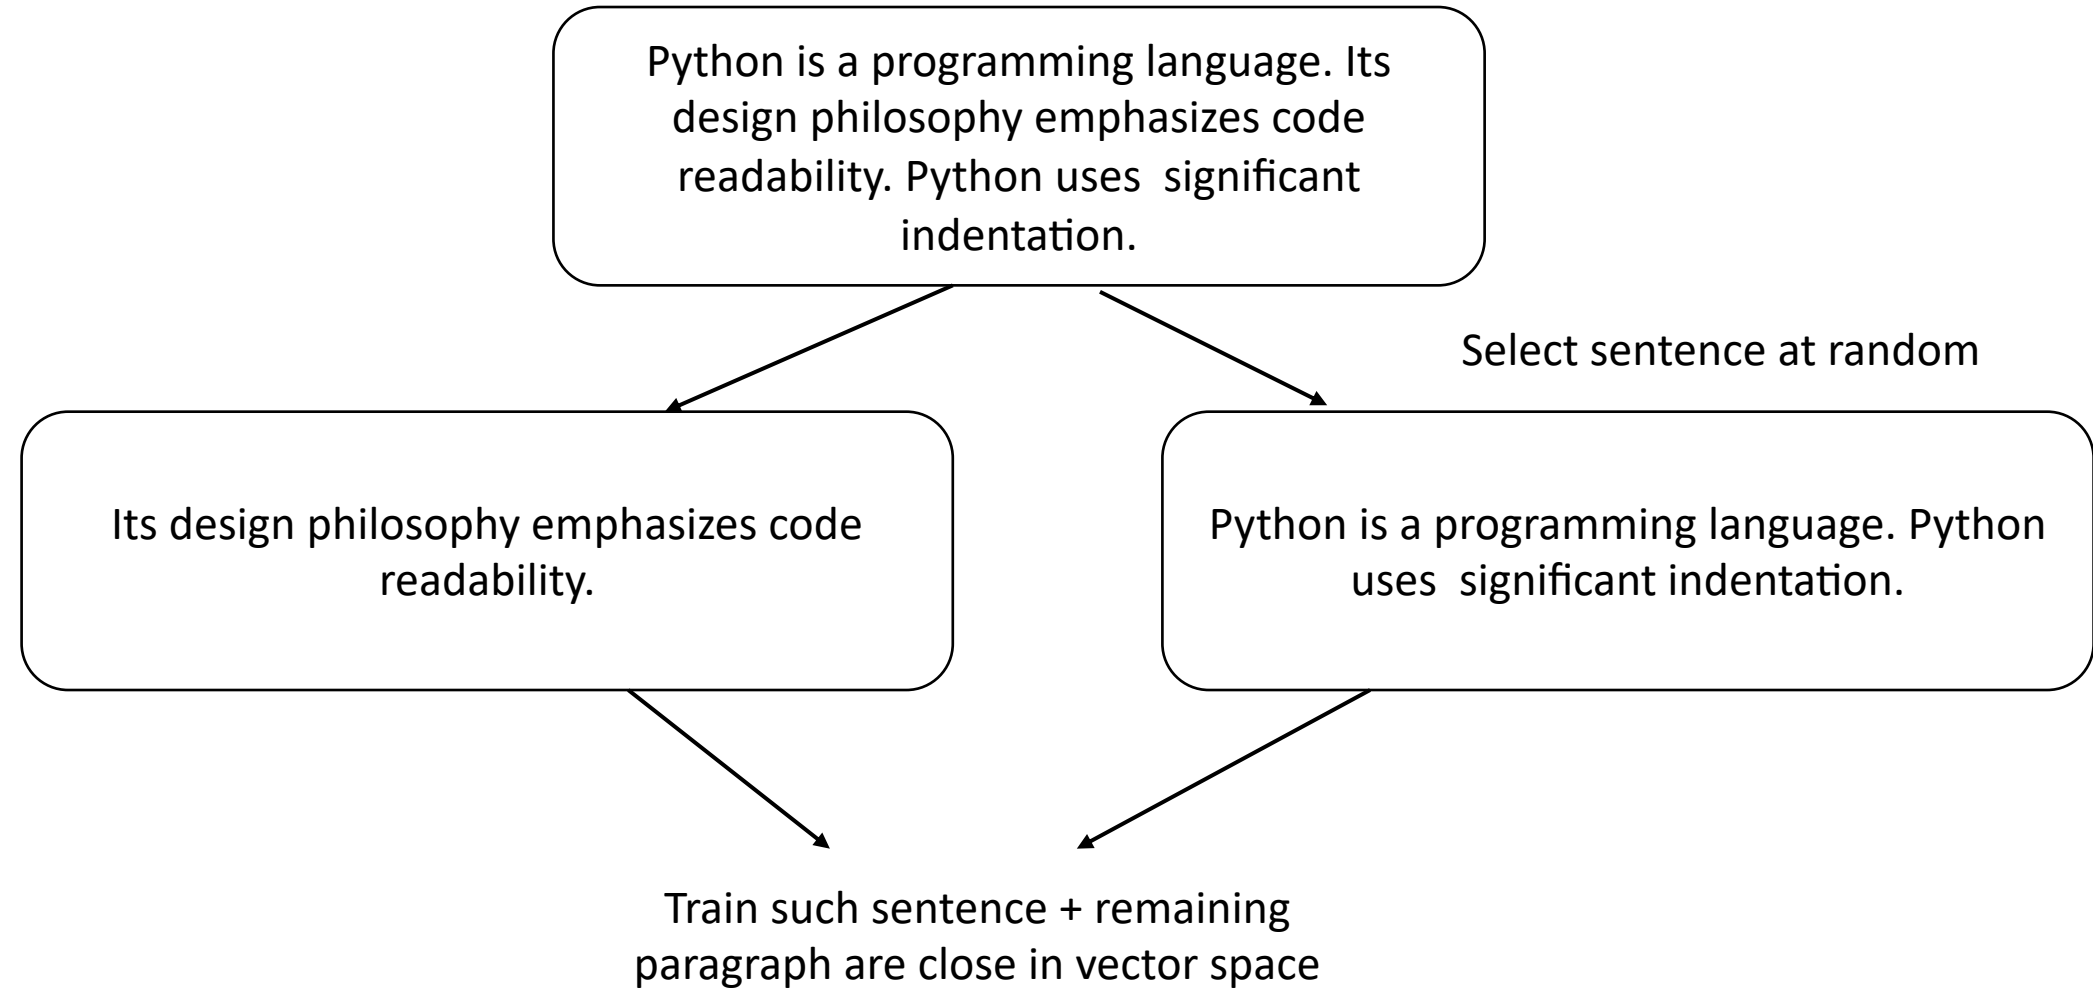


Methods for Pre-Training	Does it work?
Masked Language Modeling (MLM)	Yes
<a href="#">TSDAE</a>	Yes
<a href="#">Inverse Cloze Task</a> (ICT)	Yes
<a href="#">SimCSE</a>	No – weaker than base model
<a href="#">Contrastive Tension</a> (CT)	No – weaker than base model
<a href="#">Condenser</a> (CD)	No – weaker than base model

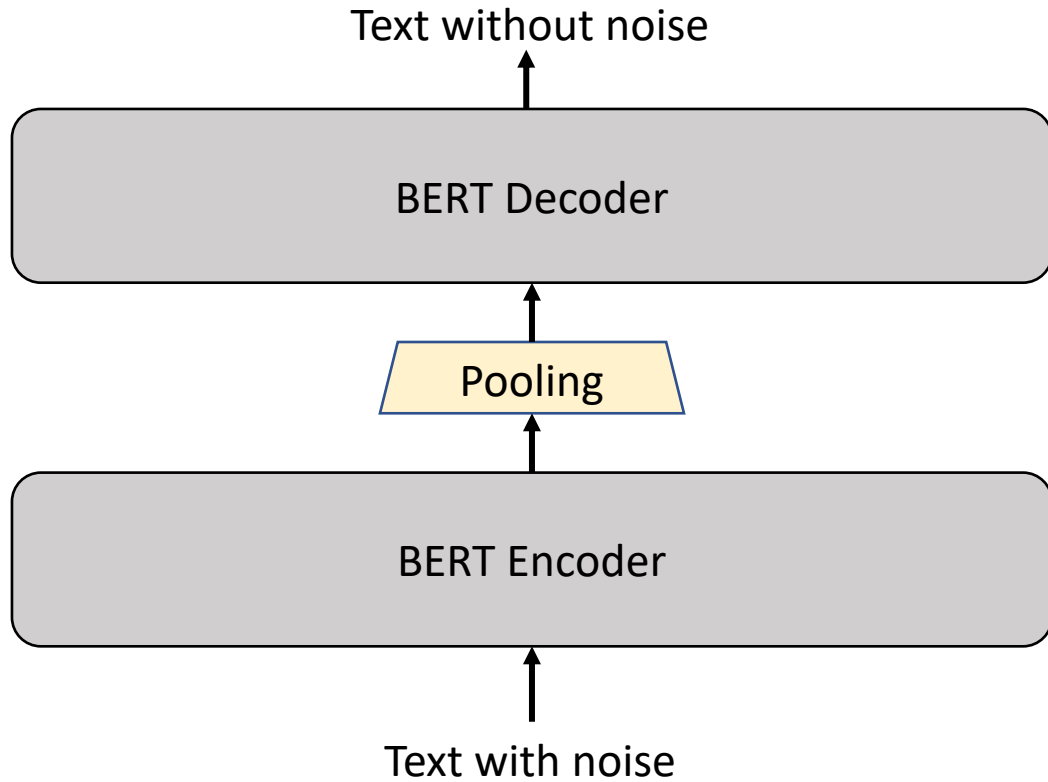
# Masked Language Model (MLM)



# Inverse Cloze Task (ICT)



# TSDAE



- Delete randomly words in the text
- Pass through the encoder
- Apply pooling to get fixed-sized text embedding
- Decoder must reconstruct text without noise from this text embedding

# Adaptive Pre-Training - Results

Models	4 Sentence Tasks	6 Dense IR Tasks
Out-of-the-box	52.3	45.2
<b>Source -&gt; Target</b>		
TSDAE	54.2	-
MLM	51.1	-
<b>Target -&gt; Source</b>		
TSDAE	56.5	49.2
MLM	55.9	46.7
ICT	-	46.5
SimCSE	52.4	45.0
CD	-	44.7
CT	53.0	44.0

# Domain Adaptation on Pre-Trained Model


✗ Adaptive pre-training is expensive

- 1) Unsupervised training on target domain
- 2) Fine-tuning on labeled source dataset (can be as large as 1 B+ training pairs)


✓ What we want:

- 1) Fine-tuning on labeled source dataset (can be as large as 1 B+ training pairs)
- 2) Unsupervised training on target domain

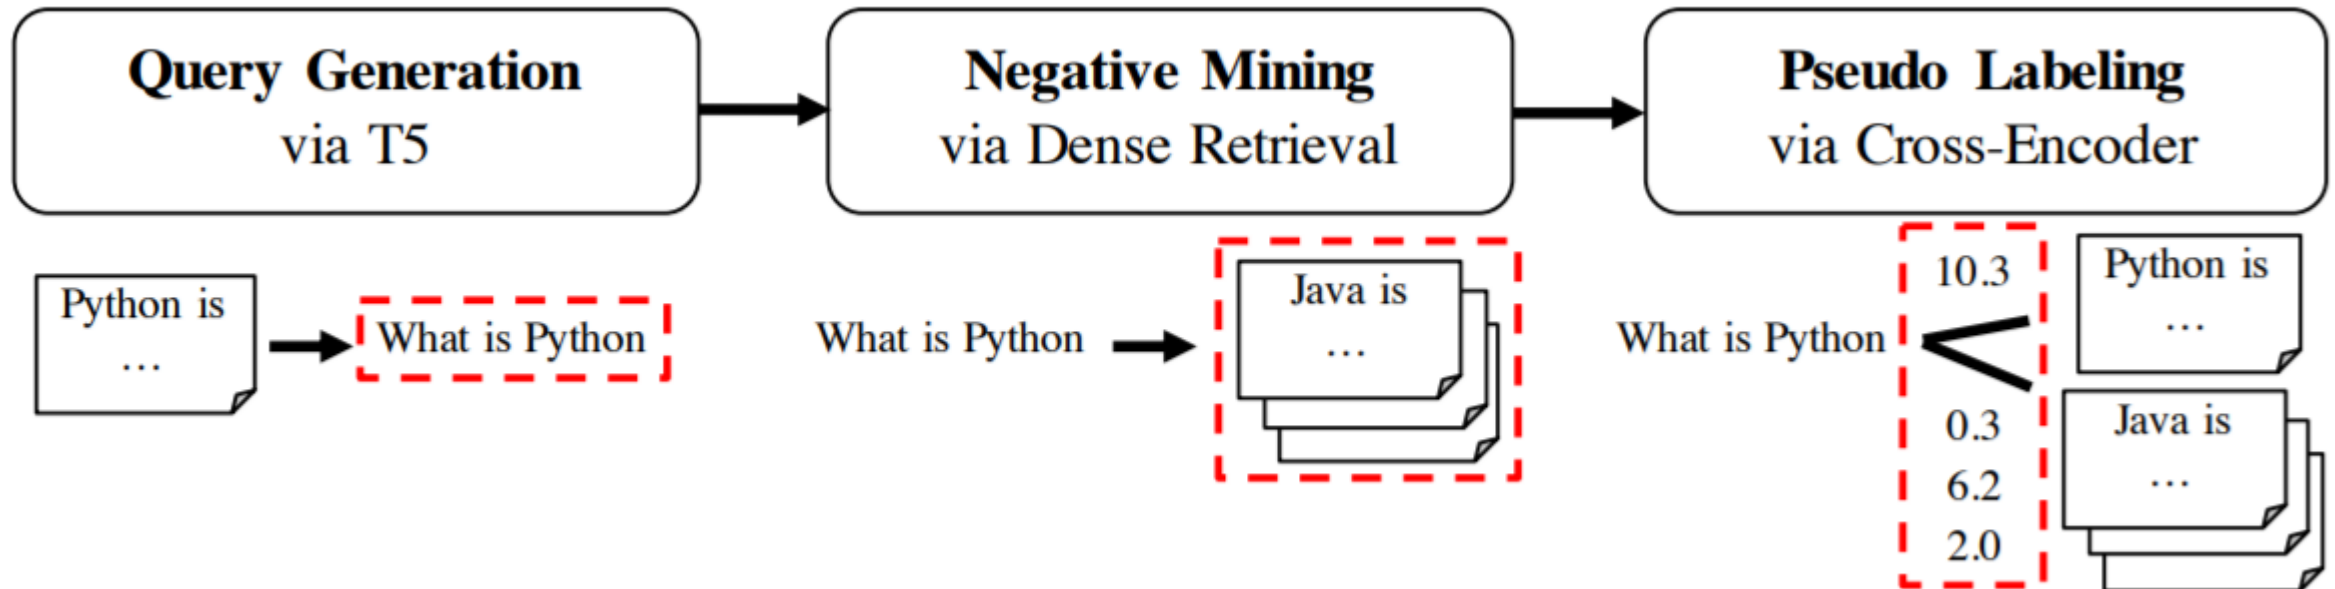
Generative Pseudo Labels (GPL) is able to achieve this



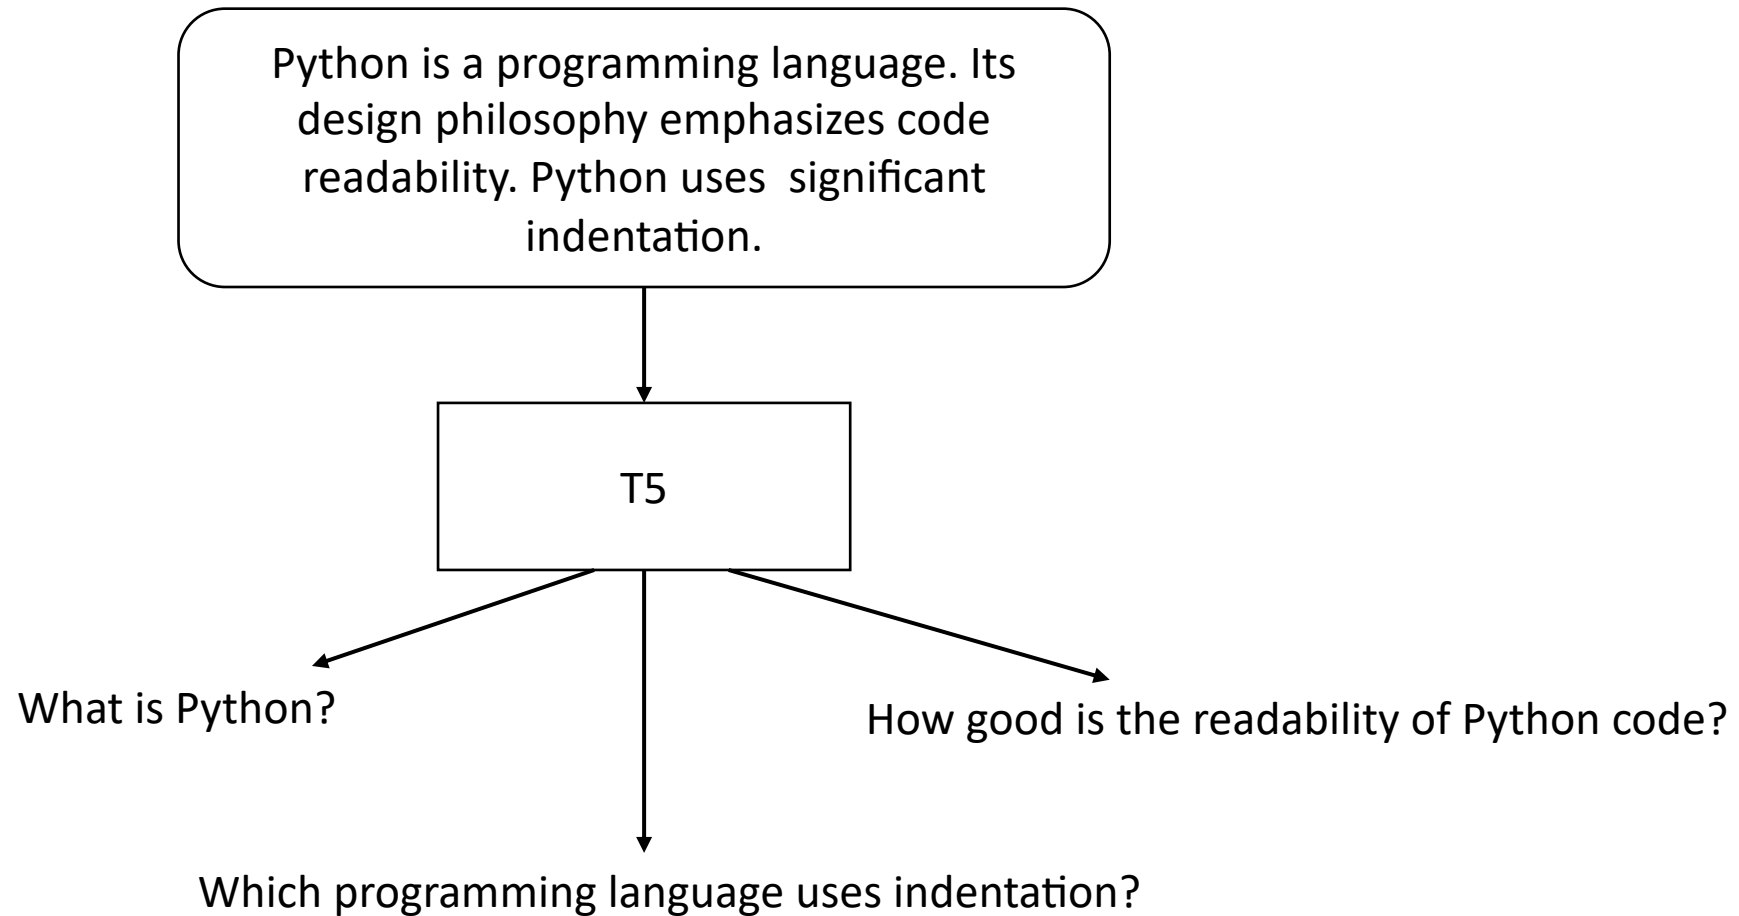
# GPL – Generative Pseudo Labeling



GPL:

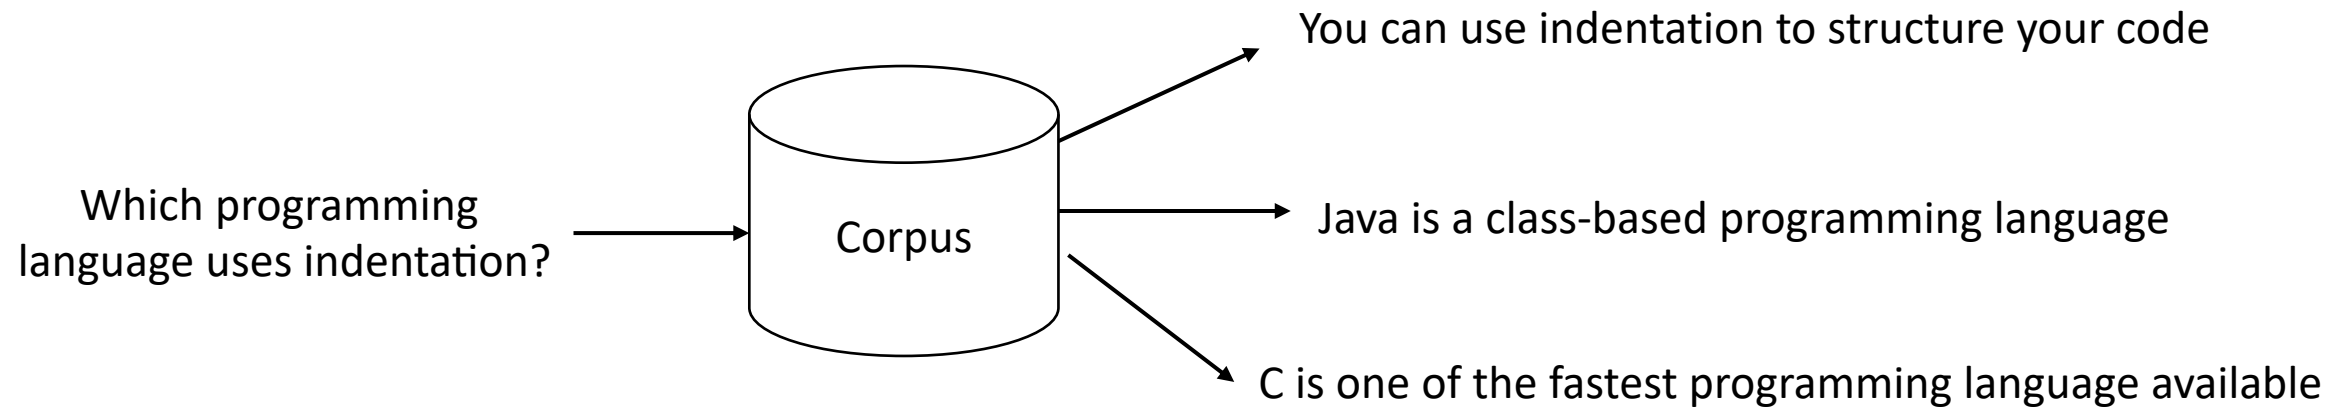


# Step 1: Generate Queries

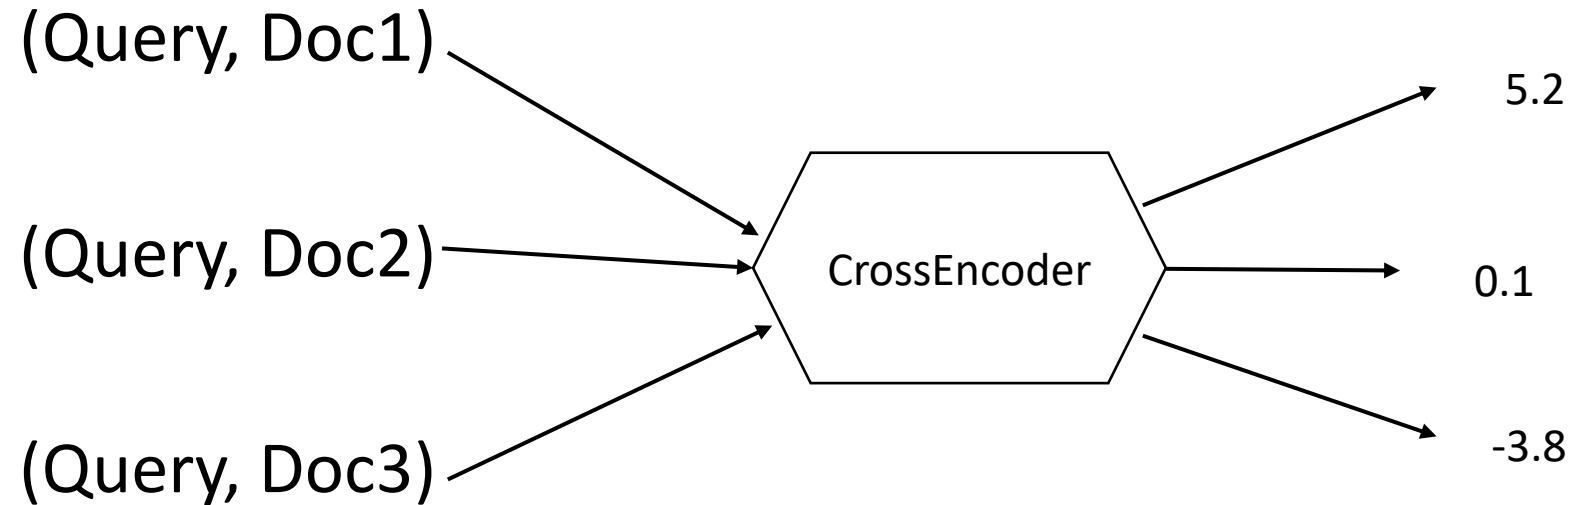




# Step 2: Mine Negatives



# Step 3: Score Pairs with CrossEncoder



# Why do we need the CrossEncoder?

	Item	Text	GPL	QGen
👉 <b>Query</b> asks for <b>definition</b> of "futures contract"	Query	what is <b>futures contract</b>	–	–
	Positive	<b>Futures contracts</b> are a member of a larger class of financial assets called derivatives ...	10.3	1
👉 <b>Easy</b> negatives: Mention "futures contract" only	Negative 1	... Anyway in this one example the s&p 500 <b>futures contract</b> has an "initial margin" of \$19,250, meaning ...	2.0	0
	Negative 2	... but the moment you exercise you must have \$5,940 in a margin account to actually use the <b>futures contract</b> ...	0.3	0
👉 <b>False</b> negative	Negative 3	... a <b>futures contract</b> is simply a contract that requires party A to buy a given amount of a commodity from party B at a specified price...	8.2	0
👉 <b>Hard</b> negative: Give partial definition	Negative 4	... A <b>futures contract</b> commits two parties to a buy/sell of the underlying securities, but ...	6.9	0

# Train Bi-Encoder with MarginMSE-Loss

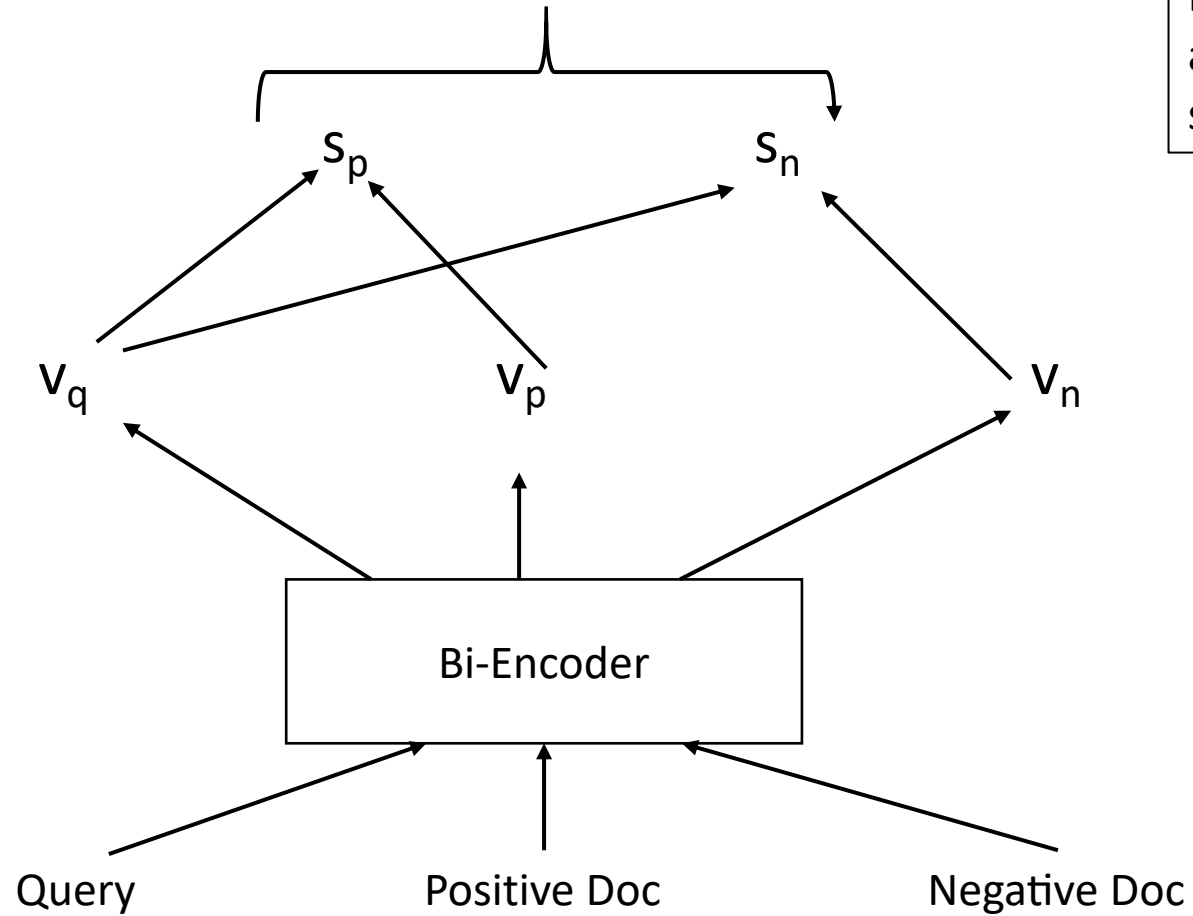
Compute Loss

$$|s_p - s_n| \text{ vs } |ce_p - ce_n|$$

CrossEncoder teaches  
BiEncoder how far vectors  
are supposed to be in vector  
space


Compute dot-scores

Compute Embeddings



# Results

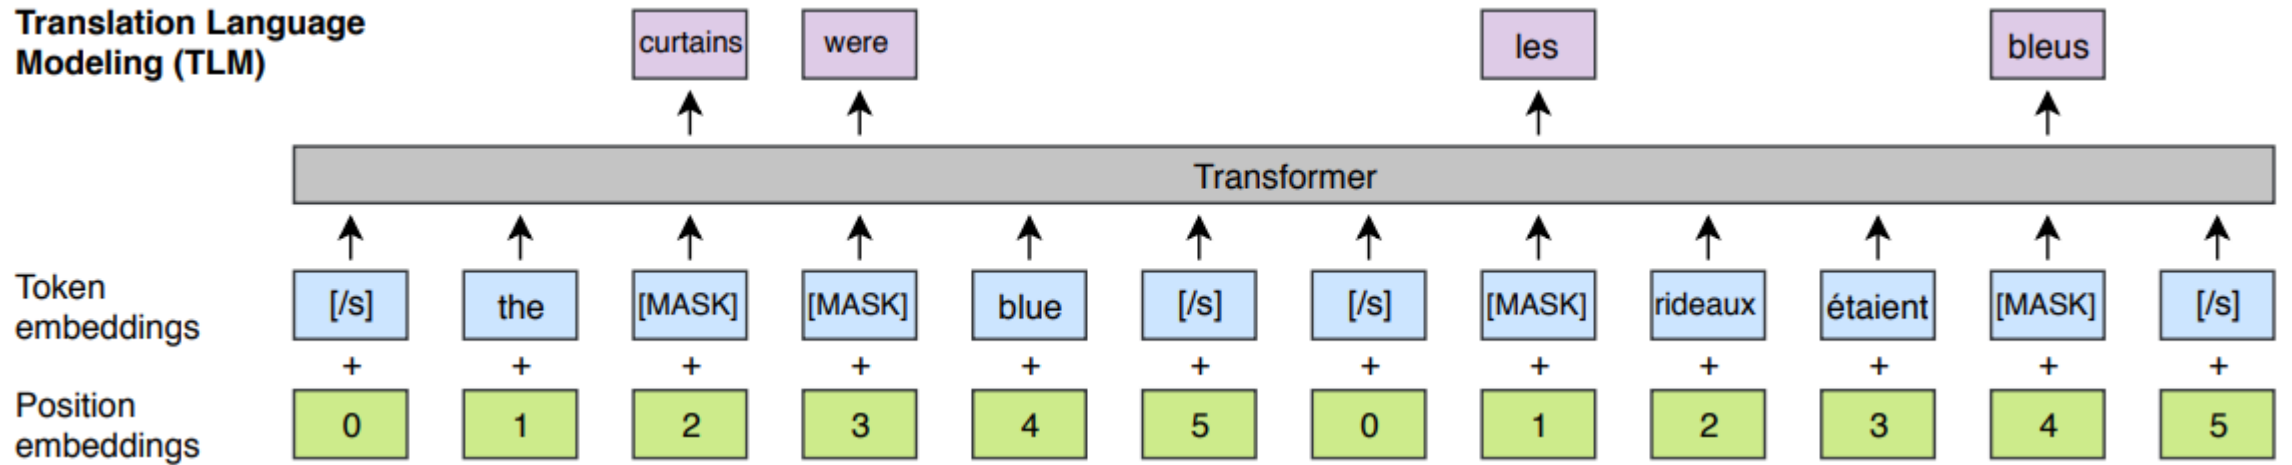
Models	6 Dense IR Tasks
Out-of-the-box	45.2
<b>Target -&gt; Source</b>	
TSDAE	49.2
MLM	46.7
<b>Generative Pseudo Labeling</b>	
GPL	51.4
TSDAE+GPL	52.4



# Multilingual Models

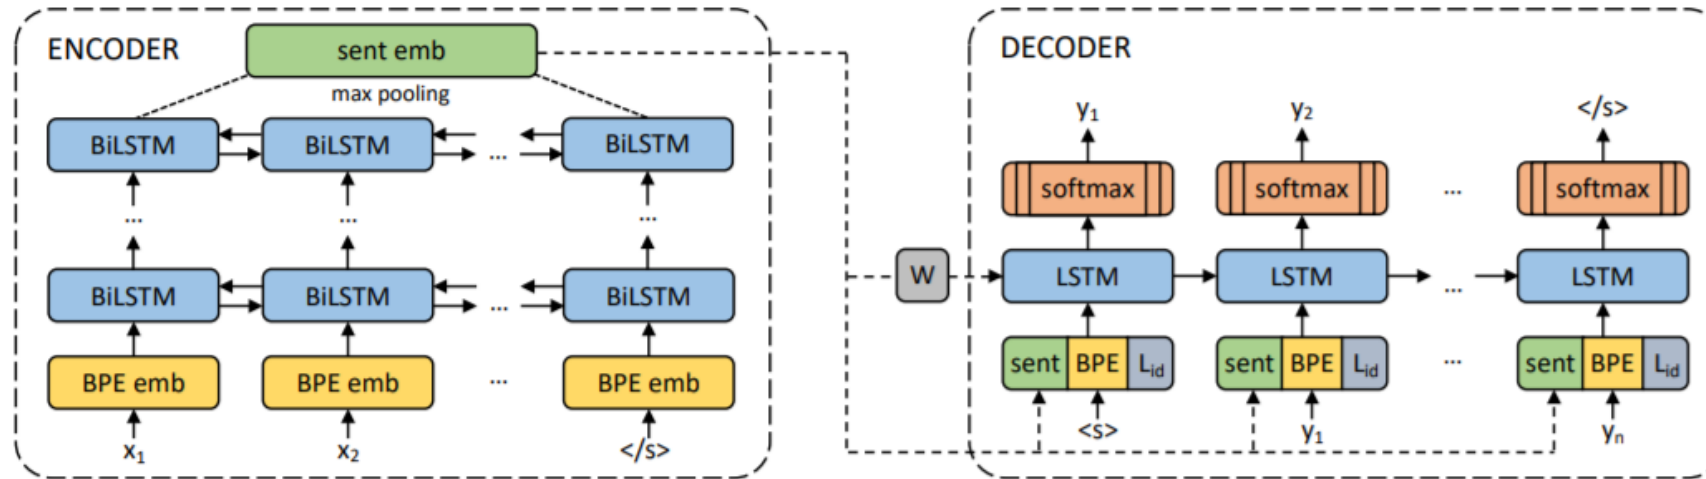
# Translation Language Model

Translation Language Modeling (TLM)



- Concatenate parallel data and run MLM

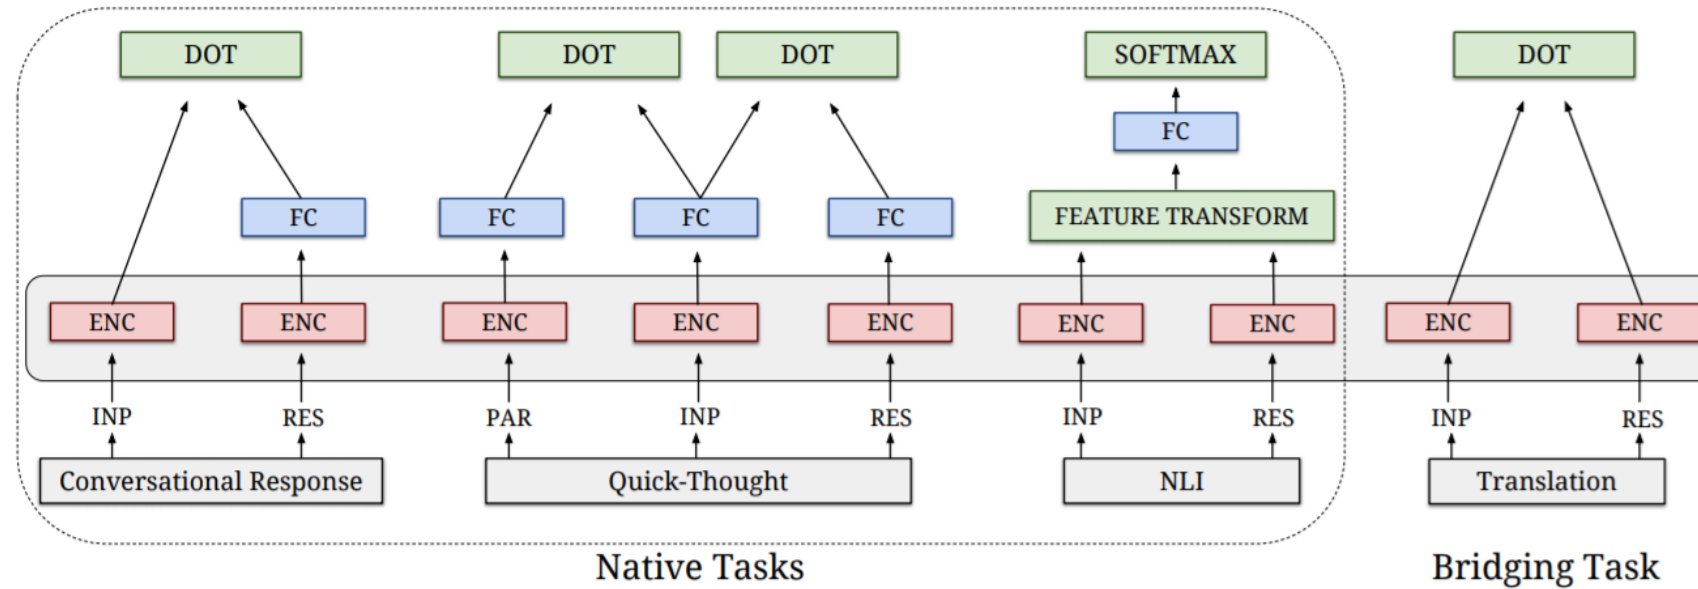
# Previous Approaches: LASER



- Use output of encoder from translation system
- Issues:
  - Cannot control what type of embeddings are learned
  - Works poorly on identifying similar sentences



# Previous Approaches: Multilingual USE



- Multi-task setup with bridging task
- Issues:
  - Getting bridging task right is challenging + requires large batch sizes
  - Hard to extend model afterwards to new languages

# LaBSE


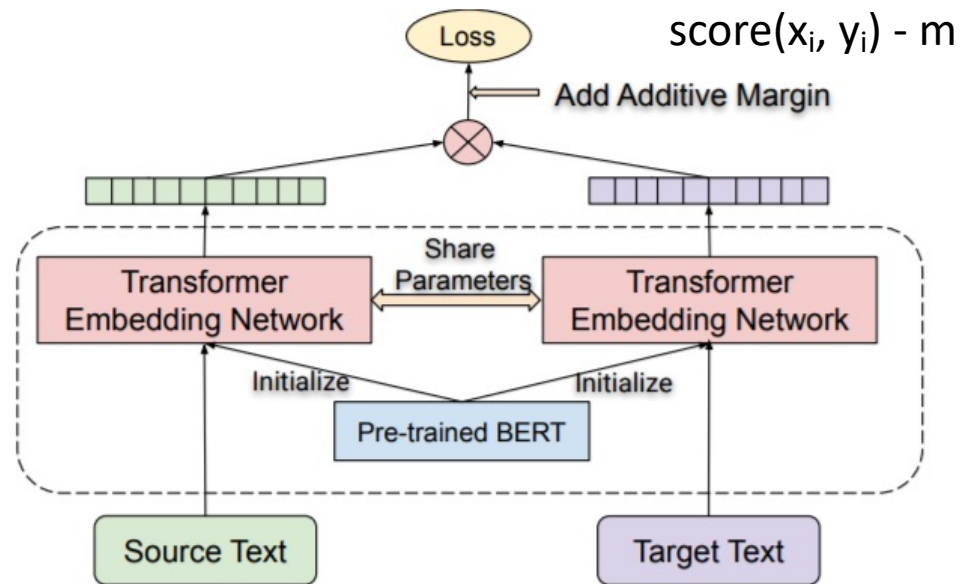
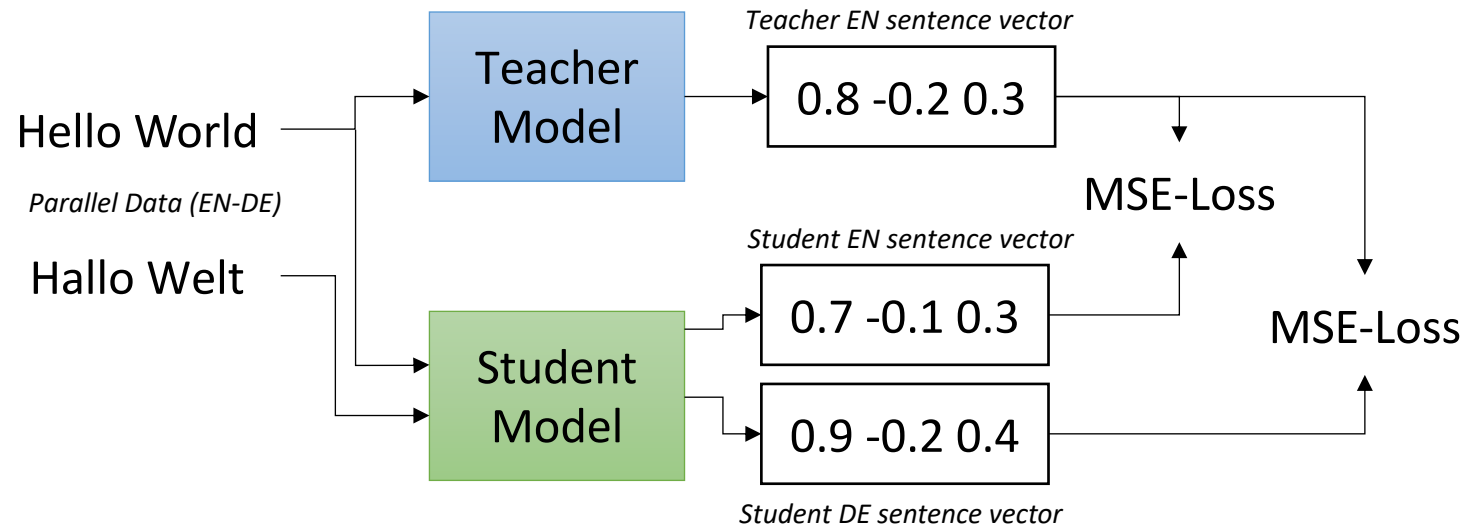


Figure 4: Average P@1 (%) on UN retrieval task of models trained with different margin values.

- Pre-Training
  - Trained on large mono-lingual dataset via MLM
  - Trained on translation pairs via TLM (Translation Lang. Model)
- Fine-tuned on translation pairs via MultipleNegativesRankingLoss

# Multilingual Knowledge Distillation



- Given:
  - Teacher sentence embedding model  $T$  (e.g. SBERT trained on English STS)
  - Parallel sentence data  $((s_1, t_1), \dots, (s_n, t_n))$
  - Student model  $S$  with multilingual vocabulary (e.g. XLM-R + Mean Pooling)
- Train student  $S$  such that:

$$S(s_i) \approx T(s_i)$$


$$S(t_i) \approx T(s_i)$$

# Performance

Model	STS	BUCC Bitext Mining
Knowledge Distillation	<b>83.7</b>	88.6
mUSE	81.1	87.7
LaBSE	73.5	<b>93.5</b>
LASER	67.0	93.0

- Models strong on multilingual STS are weak on Bitext Mining
  - Knowledge Distillation / mUSE puts similar sentences close, but which are not perfect translations

# Language Bias



- Preference of certain language combinations
- Language bias impacts performance negatively on multilingual pools
- LASER and LaBSE with strong language bias

Model	Expected Score	Actual Score	Difference
LASER	69.5	68.6	-0.92
mUSE	81.7	81.6	-0.19
LaBSE	74.4	73.1	-1.29
XLM-R ← SBERT-paraphrases	84.0	83.9	-0.11

# Language Bias – Good or Bad?

Side-effects **with** language bias:

- Same language results are ranked higher just because of language
- There might be better hits / answers in other languages

# Side-Effects without Language Bias

wedding



Who is the president?

A: Joe Biden is the current president

शादी (hindi: wedding)

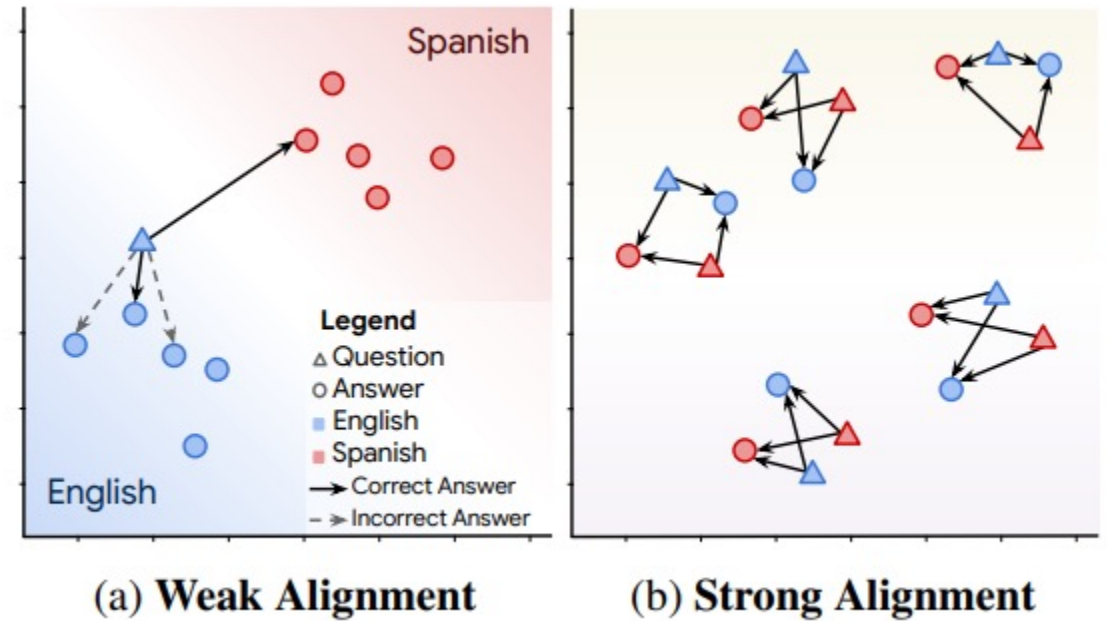


qui est le président?

A: Joe Biden is the current president

# Multilingual Search Models

- Should same language results be preferred?
  - Yes: Language Bias (weak alignment)
  - No: Strong alignment





# Batch Strategy

Q lang	A lang
en	en
en	en
en	en

(a) En-En

Q lang	A lang
ar	ar
en	en
es	es

(b) X-X

		Q lang	A lang
	Q lang	hi	hi
Q lang	en	hi	hi
ar	en	hi	hi
ar	en	en	
ar	ar		

(c) X-X-mono

Q lang	A lang
ar	en
es	hi
th	ar

(d) X-Y

- X-X is a bad idea
  - Easy to find the correct answer (just check for language)
- X-X-mono best when Q & A in same language
- X-Y best when Q & A can be in different languages

Figure 3: Sample batches for each baseline.

# Conclusion

- Training for bitext mining models:
  - LaBSE
- Training for **cross-lingual search** model
  - X-Y batch strategy
  - Getting large scale cross-lingual data is difficult

Q lang	A lang
ar	en
es	hi
th	ar

(d) X-Y



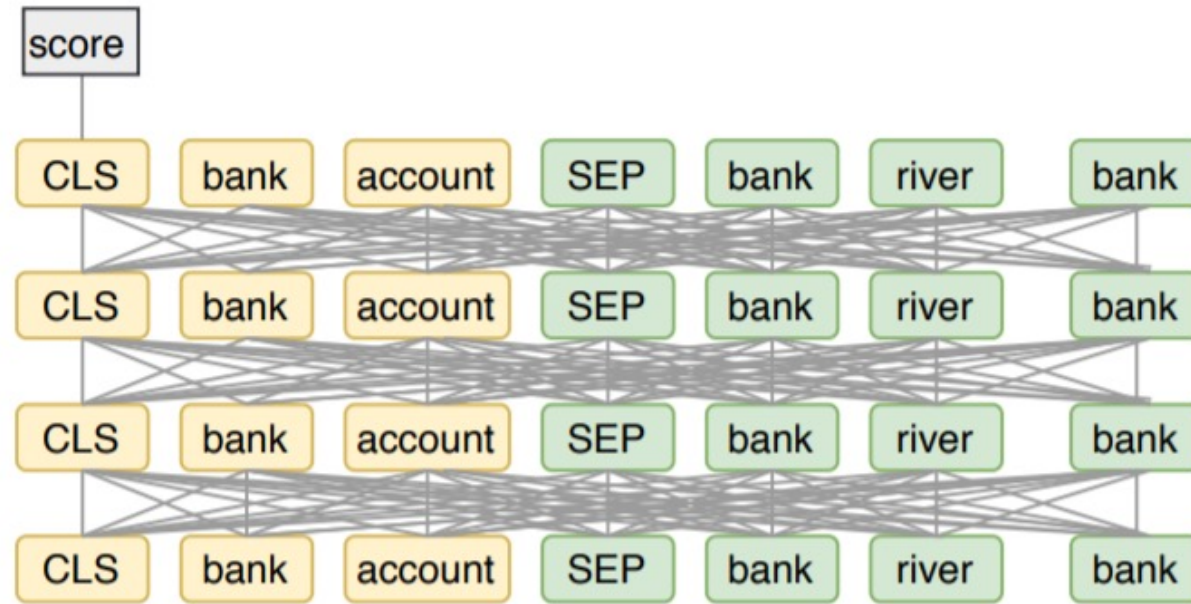
- Training for **multilingual search** models
  - X-X-mono batch strategy
  - Train on (multilingual) pair & triplet datasets
  - Add parallel data for alignment

		Q lang	A lang
		Q lang	
	Q lang	en	hi
Q lang	en	hi	hi
ar	en	hi	hi
ar	en	en	
ar	ar		

(c) X-X-mono

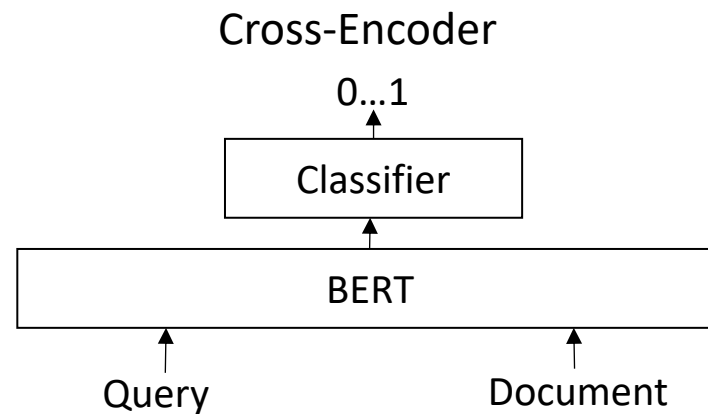
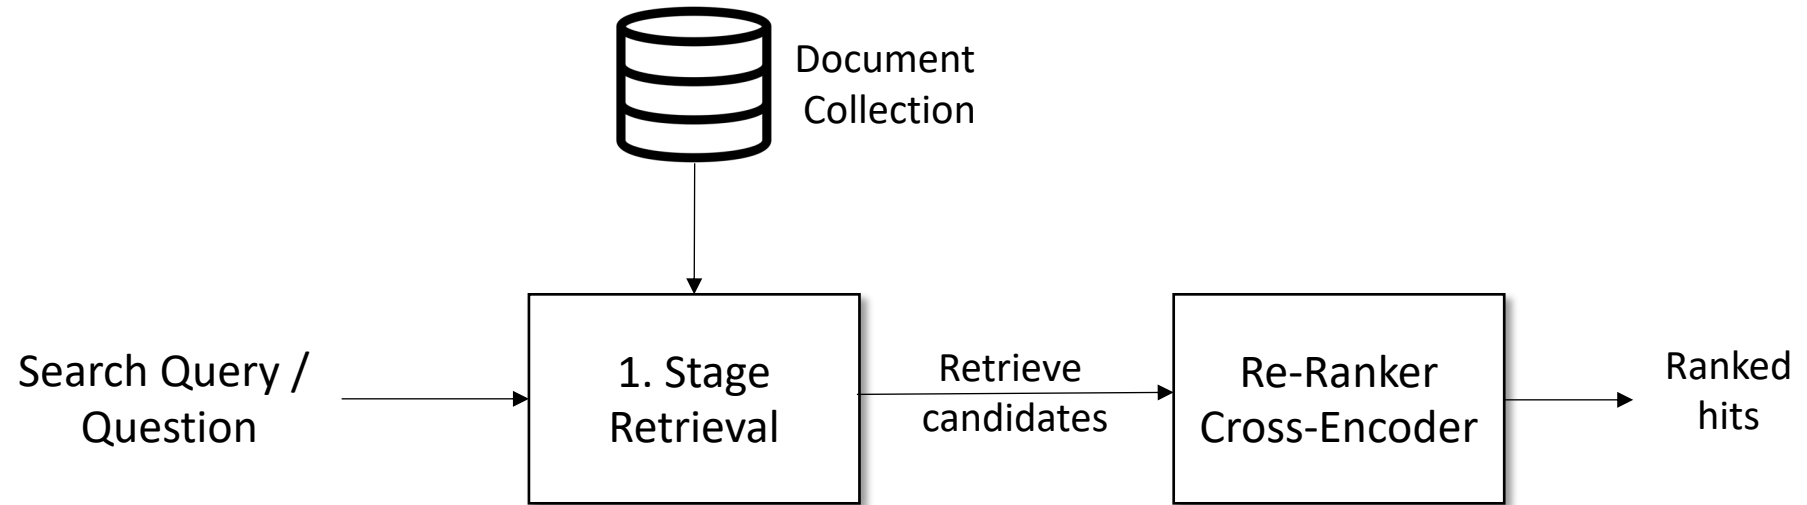
Cross-Encoder

# Cross-Encoder



- Concatenate: *Query [SEP] Passage*
- Map CLS-token output to single score

# Cross-Encoder




# Learning to Rank – Pointwise Loss

- Pointwise-Loss

- Given (query, document, label) triplets
- Set label=0 / label=1 for non- / relevant docs
- Binary classification task:  $\text{BCELoss}(\text{CE}(\text{Query}, \text{Doc}), \text{Label})$


- Challenges:

- How many non-relevant to relevant docs in the training?
- Relevance is not black & white



# Learning to Rank – Pairwise Loss

- Given (query, doc1, doc2) triplets
- Is doc1 or doc2 more relevant to the query?
- For simplification: Assume doc1 is more relevant than doc2
- RankNet Loss:
  - Compute scores:  $s^+ = \text{CE}(\text{query}, \text{doc}^+)$ ,  $s^- = \text{CE}(\text{query}, \text{doc}^-)$
  - $\text{Loss}(\text{query}, \text{doc}^+, \text{doc}^-) = \text{BCELoss}(s^+ - s^-, 1) = \log(\text{sigmoid}(s^+ - s^-))$
  - We try to maximize the margin between  $s^+$  and  $s^-$
- We don't need absolute relevance labels, just relative preferences (A or B)
- Works nice with click logs / transaction logs: Given query, what was clicked



# Learning to Rank – Listwise Loss

- Given (query, doc<sub>1</sub>, doc<sub>2</sub>, doc<sub>3</sub>, ...)
- Which doc is the most relevant for the query?
- Many loss functions available: LambdaRank, LambdaMART, ApproxNDCG, NeuralNDCG...
  - Often they try to optimize the eval measure (like nDCG)
  - I didn't observe large differences
- I prefer ListRank Loss / ListNet Loss:
  - Compute  $s_1 = \text{CE}(\text{query}, \text{doc}_1)$ ,  $s_2 = \text{CE}(\text{query}, \text{doc}_2)$ ,  $s_3 = \text{CE}(\text{query}, \text{doc}_3)$ , ...
  - `CrossEntropyLoss( [s1, s2, s3, ...], label )`
  - Label: Which document is the most relevant?
  - Train with 1 positive and many negative docs
  - With 1 negative: Identical to Pairwise Loss / RankNet Loss




# Learning to Rank

- Pointwise loss performs the worse
  - Hard to tell what docs are relevant / irrelevant
  - Hard to select the ratio of positive vs negative labels
  - Harder to get labeled data
- Pairwise / Listwise Loss performs better
  - Just relative importance is relevant (is A or B better?)
  - Easier to extract from click logs / transaction logs

# Importance of Negatives

- Listwise loss: [query, positive, neg1, neg2, ...]
- Negatives are either from top-50 or top-1k from BM25




# Importance of Negatives

Train Neg↓ / Inference→	BM25	BM25*	HDCT
BM25	39.6	39.5	38.1
BM25*	40.7	<b>42.3</b>	41.8
HDCT	40.8	41.9	<b>43.4</b>

- Performance drops if train sample is different from test first-stage retrieval system
- As we optimize for unknown first-stage system:
  - Samples negatives from different systems (lexical & embedding based)

# Number of Negatives



# Multilingual Cross-Encoder

- Trained on Machine-Translated MS MARCO (incl. de, ar, id, ru)
- Performance on GermanQuAD & Mr. Tydi (Arabic, Indonesian, Russian)

Model	Performance
mdeberta v3 (training: English only)	52.2
mdeberta v3 (14 mMARCO langs)	53.0
LaBSE	52.6
mMiniLM	52.0

- Models perform surprisingly well even when train on English only